




Understanding the Lomb–Scargle Periodogram

Jacob T. VanderPlas 

University of Washington, eScience Institute, 3910 15th Ave NE, Seattle WA 98195

Received 2017 August 26; revised 2017 December 6; accepted 2017 December 7; published 2018 May 11

Abstract

The Lomb–Scargle periodogram is a well-known algorithm for detecting and characterizing periodic signals in unevenly sampled data. This paper presents a conceptual introduction to the Lomb–Scargle periodogram and important practical considerations for its use. Rather than a rigorous mathematical treatment, the goal of this paper is to build intuition about what assumptions are implicit in the use of the Lomb–Scargle periodogram and related estimators of periodicity, so as to motivate important practical considerations required in its proper application and interpretation.

Key words: methods: data analysis – methods: statistical

1. Introduction

The Lomb–Scargle periodogram (Lomb 1976; Scargle 1982) is a well-known algorithm for detecting and characterizing periodicity in unevenly sampled time-series and has seen particularly wide use within the astronomy community. As an example of a typical application of this method, consider the data shown in Figure 1: this is an irregularly sampled time-series showing a single object from the LINEAR survey (Sesar et al. 2011; Palaversa et al. 2013), with unfiltered magnitude measured 280 times over the course of 5.5 yr. By eye, it is clear that the brightness of the object varies in time with a range spanning approximately 0.8 mag, but what is not immediately clear is that this variation is periodic in time. The Lomb–Scargle periodogram is a method that allows efficient computation of a Fourier-like power spectrum estimator from such unevenly sampled data, resulting in an intuitive means of determining the period of oscillation.

The Lomb–Scargle periodogram computed from these data is shown in the left panel of Figure 2. The Lomb–Scargle periodogram here yields an estimate of the Fourier power as a function of period of oscillation, from which we can read off the period of oscillation of approximately 2.58 hr. The right panel of Figure 2 shows a folded visualization of the same data as Figure 1—i.e., plotted as a function of phase rather than time.

Often this is exactly how the Lomb–Scargle periodogram is presented: as a clean, well-defined procedure to generate a power spectrum and to detect the periodic component in an unevenly sampled data set. In practice, however, there are a number of subtle issues that must be considered when applying a Lomb–Scargle analysis to real-world data sets. Here are a few questions in particular that we might wish to ask about the results in Figure 2:

1. How does the Lomb–Scargle periodogram relate to the classical Fourier power spectrum?
2. What is the source of the pattern of multiple peaks revealed by the Lomb–Scargle periodogram?
3. What is the largest frequency (i.e., Nyquist-like limit) that such an analysis is sensitive to?
4. How should we choose the spacing of the frequency grid for our periodogram?
5. What assumptions, if any, does the Lomb–Scargle approach make regarding the form of the unknown signal?

6. How should we understand and report the uncertainty of the determined frequency?

Quantitative treatments of these sorts of questions are presented in various textbooks and review papers, but I have not come across any single concise reference that gives a good intuition for how to think about such questions. This paper seeks to fill that gap and provide a practical, just-technical-enough guide to the effective use of the Lomb–Scargle method for detection and characterization of periodic signals. This paper does not seek a complete or rigorous treatment of the mathematics involved, but rather seeks to develop the intuition of *how* to think about these questions, with references to relevant technical treatments.

1.1. Why Lomb–Scargle?

Before we begin exploring the Lomb–Scargle periodogram in more depth, it is worth briefly considering the broader context of methods for detecting and characterizing periodicity in time-series. First, it is important to note that there are many different modes of time-series observation. Point observation like that shown in Figure 1 is typical of optical astronomy: values (often with uncertainties) are measured at discrete points in time, which may be equally or unequally spaced. Other modes of observation—e.g., time-tag events, binned event data, time-to-spill events, etc.—are common in high-energy astronomy and other areas. We will not consider such event-driven data modes here, but note that there have been some interesting explorations of unified statistical treatments of all of the above modes of observation (e.g., Scargle 1998, 2002).

Even limiting our focus to point observations, there are a large number of complementary techniques for periodic analysis, which generally can be categorized into a few broad categories:

Fourier methods are based on the Fourier transform, power spectra, and closely related correlation functions. These methods include the classical or Schuster periodogram (Schuster 1898), the Lomb–Scargle periodogram (Lomb 1976; Scargle 1982), the correlation-based method of Edelson & Krolik (1988), and related approaches (see also Foster 1996, for a discussion of wavelet transforms in this context).




Figure 1. Observed light curve from LINEAR object ID 11375941. Uncertainties are indicated by the gray error bars on each point. The Python code to reproduce this figure, as well as all other figures in this manuscript, is available at <http://github.com/jakevdp/PracticalLombScargle/>.

Phase-folding methods depend on folding observations as a function of phase, computing a cost function across the phased data (often within bins constructed across phase space) and optimizing this cost function across candidate frequencies. Some examples are String Length (Dworetzky 1983), Analysis of Variance (Schwarzenberg-Czerny 1989), Phase Dispersion Minimization (Stellingwerf 1978), the Gregory–Laredo method (Gregory & Laredo 1992), and the conditional entropy method (Graham et al. 2013b). Methods based on correntropy are similar in spirit but do not always require explicit phase folding (Huijse et al. 2011, 2012).

Least-squares methods involve fitting a model to the data at each candidate frequency and selecting the frequency that maximizes the likelihood. The Lomb–Scargle periodogram also falls in this category (see Section 5), as does the Supersmoothing approach (Reimann 1994). Other studies recommend statistics other than least-squares residuals; see, e.g., the orthogonal polynomial fits of Schwarzenberg-Czerny (1996).

Bayesian approaches apply Bayesian probability theory to the problem, often in a similar manner to the phase-folding and/or least-squares approaches. Examples are the generalized Lomb–Scargle models of Bretthorst (1988), the phase-binning model of Gregory & Laredo (1992), Gaussian process models (e.g., Wang et al. 2012), and models based on stochastic processes (e.g., Kelly et al. 2014).

Various reviews have been undertaken to compare the efficiency and effectiveness of the available methods; for example, Schwarzenberg-Czerny (1999) focuses on the statistical properties of methods and recommends those based on smooth model fits over methods based on phase binning, while Graham et al. (2013a) instead take an empirical approach and find that when considering detection efficiency in real data sets, no suitably efficient algorithm universally outperforms the others.

In light of this breadth of available methods, why limit our focus here to the Lomb–Scargle periodogram? One reason is cultural: the Lomb–Scargle periodogram is perhaps the best-known technique to compute periodicity of unequally spaced data in astronomy and other fields, and so it is the first tool many will reach for when searching for periodic content in a signal. But there is a deeper reason as well; it turns out that Lomb–Scargle occupies a unique niche: it is motivated by

Fourier analysis (see Section 5), but it can also be viewed as a least-squares method (see Section 6). It can be derived from the principles of Bayesian probability theory (see Section 6.5) and has been shown to be closely related to bin-based phase-folding techniques under some circumstances (see Swingler 1989). Thus, the Lomb–Scargle periodogram occupies a unique point of correspondence between many classes of methods and so provides a focus for discussion that has bearing on considerations involved in *all* of these methods.

1.2. Outline

The remainder of this paper is organized as follows:

- Section 2* presents a review of the continuous Fourier transform and some of its useful properties, including defining the notion of a power spectrum (i.e., classical/Schuster periodogram) for detecting periodic content in a signal.
- Section 3* builds on these properties to explore how different observation patterns (i.e., window functions) affect the periodogram and discusses a conceptual understanding of the Nyquist limiting frequency.
- Section 4* considers nonuniformly sampled signals as a special case of an observational window and shows that the Nyquist-like limit for this case is quite different than what many practitioners in the field often assume.
- Section 5* introduces the Lomb–Scargle periodogram, a modified version of the classical periodogram for unevenly sampled data, as well as the motivation behind these modifications.
- Section 6* discusses a complementary view of the Lomb–Scargle periodogram as the result of least-squares model fitting, as well as some extensions enabled by that viewpoint.
- Section 7* builds on the concepts introduced in earlier sections to discuss important practical considerations for the use of the Lomb–Scargle periodogram, including the choice of frequency grid, uncertainties and false-alarm probabilities (FAPs), and modes of failure that should be accounted for when working with real-world data.
- Section 8* concludes and summarizes the key recommendations for practical use of the Lomb–Scargle periodogram.

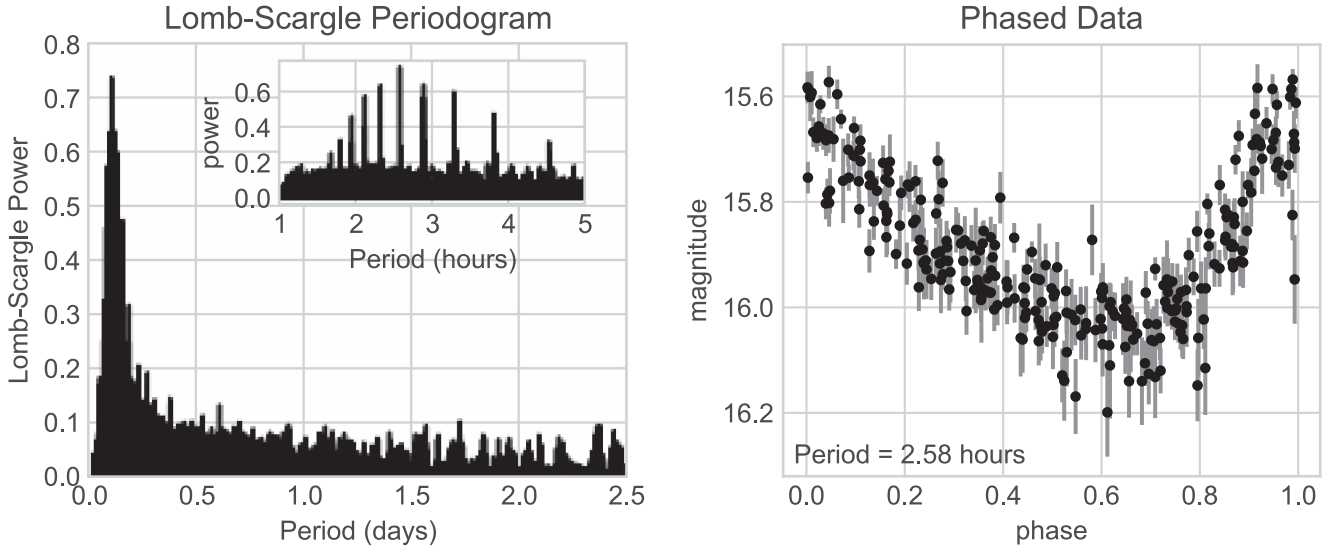


Figure 2. Left panel: Lomb–Scargle periodogram computed from the data shown in Figure 1, with an inset detailing the region near the peak. Right panel: input data in Figure 1, folded over the detected 2.58 hr period to show the coherent periodic variability. For more discussion of this particular object, see Palaversa et al. (2013).

2. Background: The Continuous Fourier Transform

In order to understand how we should interpret the Lomb–Scargle periodogram, we will first briefly step back and review the subject of Fourier analysis of continuous signals. Consider a continuous signal $g(t)$. Its Fourier transform is given by the following integral, where $i \equiv \sqrt{-1}$ denotes the imaginary unit:

$$\hat{g}(f) \equiv \int_{-\infty}^{\infty} g(t) e^{-2\pi i f t} dt. \quad (1)$$

The inverse relationship is given by

$$g(t) \equiv \int_{-\infty}^{\infty} \hat{g}(f) e^{+2\pi i f t} df. \quad (2)$$

For convenience we will also define the Fourier transform operator \mathcal{F} , such that

$$\mathcal{F}\{g\} = \hat{g}, \quad (3)$$

$$\mathcal{F}^{-1}\{\hat{g}\} = g. \quad (4)$$

The functions g and \hat{g} are known as a *Fourier pair*, which we will sometimes denote as $g \longleftrightarrow \hat{g}$.

2.1. Properties of the Fourier Transform

The continuous Fourier transform has a number of useful properties that we will make use of in our discussion.

The Fourier transform is a linear operation. That is, for any constant A and any functions $f(t)$ and $g(t)$, we can write

$$\begin{aligned} \mathcal{F}\{f(t) + g(t)\} &= \mathcal{F}\{f(t)\} + \mathcal{F}\{g(t)\} \\ \mathcal{F}\{A f(t)\} &= A \mathcal{F}\{f(t)\}. \end{aligned} \quad (5)$$

Both identities follow from the linearity of the Fourier integral.

The Fourier transform of a sinusoid with frequency f_0 is a sum of delta functions at $\pm f_0$. From the integral definition of the Dirac delta function,¹ we can write

$$\mathcal{F}\{e^{2\pi i f_0 t}\} = \delta(f - f_0). \quad (6)$$

Using Euler’s formula² for the complex exponential along with the linearity of the Fourier transform leads to the following identities:

$$\begin{aligned} \mathcal{F}\{\cos(2\pi f_0 t)\} &= \frac{1}{2} [\delta(f - f_0) + \delta(f + f_0)] \\ \mathcal{F}\{\sin(2\pi f_0 t)\} &= \frac{1}{2i} [\delta(f - f_0) - \delta(f + f_0)]. \end{aligned} \quad (7)$$

In other words, a sinusoidal signal with frequency f_0 has a Fourier transform consisting of a weighted sum of delta functions at $\pm f_0$.

A time shift imparts a phase in the Fourier transform. Given a well-behaved function $g(t)$, we can use a transformation of variables to derive the following identity:

$$\mathcal{F}\{g(t - t_0)\} = \mathcal{F}\{g(t)\} e^{-2\pi i f t_0}. \quad (8)$$

Notice that the time shift does not change the amplitude of the resulting transform, but only the phase.

These properties taken together make the Fourier transform quite useful for the study of periodic signals. The linearity of the transform means that any signal made up of a sum of sinusoidal components will have a Fourier transform consisting of a sum of delta functions marking the frequencies of those sinusoids—that is, the Fourier transform directly measures additive periodic content in a continuous function.

Further, if we compute the squared amplitude of the resulting transform, we can both do away with complex components and remove the phase imparted by the choice of temporal baseline; this squared amplitude is usually known as the *power spectral density* (PSD) or simply the *power spectrum*:

$$\mathcal{P}_g \equiv |\mathcal{F}\{g\}|^2. \quad (9)$$

The power spectrum of a function is a positive real-valued function of the frequency f that quantifies the contribution of each frequency f to the total signal. Note that if g is real-valued, it follows that \mathcal{P}_g is an even function, i.e., $\mathcal{P}_g(f) = \mathcal{P}_g(-f)$.

¹ $\delta(f) \equiv \int_{-\infty}^{\infty} e^{-2\pi i x f} df$.

² $e^{ix} = \cos x + i \sin x$.




Figure 3. Visualization of important Fourier pairs.

2.2. Some Useful Fourier Pairs

We have already discussed that the Fourier transform of a complex exponential is a single delta function. This is just one of many Fourier pairs to keep in mind as we progress to understanding the Lomb–Scargle periodogram. We list a few more important pairs here (see Figure 3 for a visual representation of the following pairs):

The Fourier transform of a sinusoid is a pair of Delta functions (see Figure 3(a)):

$$\mathcal{F}\{\cos(2\pi f_0 t)\} = \frac{1}{2}[\delta(f - f_0) + \delta(f + f_0)]. \quad (10)$$

We saw this above, but we repeat it here for completeness.

The Fourier transform of a Gaussian is a Gaussian (see Figure 3(b)):

$$\mathcal{F}\{N(t; \sigma)\} = \frac{1}{\sqrt{2\pi\sigma^2}} N\left(f; \frac{1}{2\pi\sigma}\right). \quad (11)$$

The Gaussian function $N(t, \sigma)$ is given by

$$N(t; \sigma) \equiv \frac{1}{\sqrt{2\pi\sigma^2}} e^{-t^2/(2\sigma^2)}. \quad (12)$$

The Fourier transform of a rectangular function is a sinc function (see Figure 3(c)):

$$\mathcal{F}\{\Pi_T(t)\} = \text{sinc}(fT). \quad (13)$$

The rectangular function, $\Pi(t)$, is a normalized symmetric function that is uniform within a range given by T and zero elsewhere:

$$\Pi_T(t) \equiv \begin{cases} 1/T, & |t| \leq T/2 \\ 0, & |t| > T/2. \end{cases} \quad (14)$$

The sinc function is given by the standard definition:

$$\text{sinc}(x) \equiv \frac{\sin(\pi x)}{\pi x}. \quad (15)$$

The Fourier transform of a Dirac comb is a Dirac comb (see Figure 3(d)):

$$\mathcal{F}\{\text{III}_T(t)\} = \frac{1}{T} \text{III}_{1/T}(f). \quad (16)$$

The Dirac comb, $\text{III}_T(t)$, is an infinite sequence of Dirac delta functions placed at even intervals of size T :

$$\text{III}_T(t) \equiv \sum_{n=-\infty}^{\infty} \delta(t - nT). \quad (17)$$

Notice in each of these Fourier pairs the reciprocity of scales between a function and its Fourier transform: a narrow function will have a broad transform, and vice versa. More quantitatively, a function with a characteristic scale T will in general have a Fourier transform with a characteristic scale of $1/T$. Such reciprocity is central to many diverse applications of Fourier analysis, from electrodynamics to music theory to quantum mechanics. For our purposes, this property will turn out to be quite important as we push further in understanding the Lomb–Scargle periodogram.

2.3. The Convolution Theorem

A final property of the Fourier transform that we will discuss here is its ability to convert convolutions into pointwise products. A convolution of two functions, usually denoted by




Figure 4. Visualization of a convolution between a continuous signal and a rectangular smoothing kernel. The normalized rectangular window slides across the domain (top panel), such that at each point the mean of the values within the window is used to compute the smoothed function (bottom panel).




Figure 5. Visualization of the convolution theorem (Equation (19)). Recall that the Fourier transform of a convolution is the pointwise product of the two Fourier transforms. In the right panels, the black and gray lines represent the real and imaginary parts of the transform, respectively.

the $*$ symbol, is defined as follows:

$$[f * g](t) \equiv \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau. \quad (18)$$

The convolution can be thought of as an operation that “slides” one function past the other, integrating the product at each step. Such an operation is commonly used, for example, in smoothing a function, as visualized in Figure 4 for a rectangular smoothing window.

Given this definition of a convolution, it can be shown that the Fourier transform of a convolution is the pointwise product

of the individual Fourier transforms:

$$\mathcal{F}\{f * g\} = \mathcal{F}\{f\} \cdot \mathcal{F}\{g\}. \quad (19)$$

In practice, this can be a much more efficient means of numerically computing a convolution than to directly solve at each time t the integral over τ that appears in Equation (18). The identity in Equation (19) is known as the *convolution theorem* and is illustrated in Figure 5. An important corollary is that the Fourier transform of a product is a convolution of the two transforms:

$$\mathcal{F}\{f \cdot g\} = \mathcal{F}\{f\} * \mathcal{F}\{g\}. \quad (20)$$

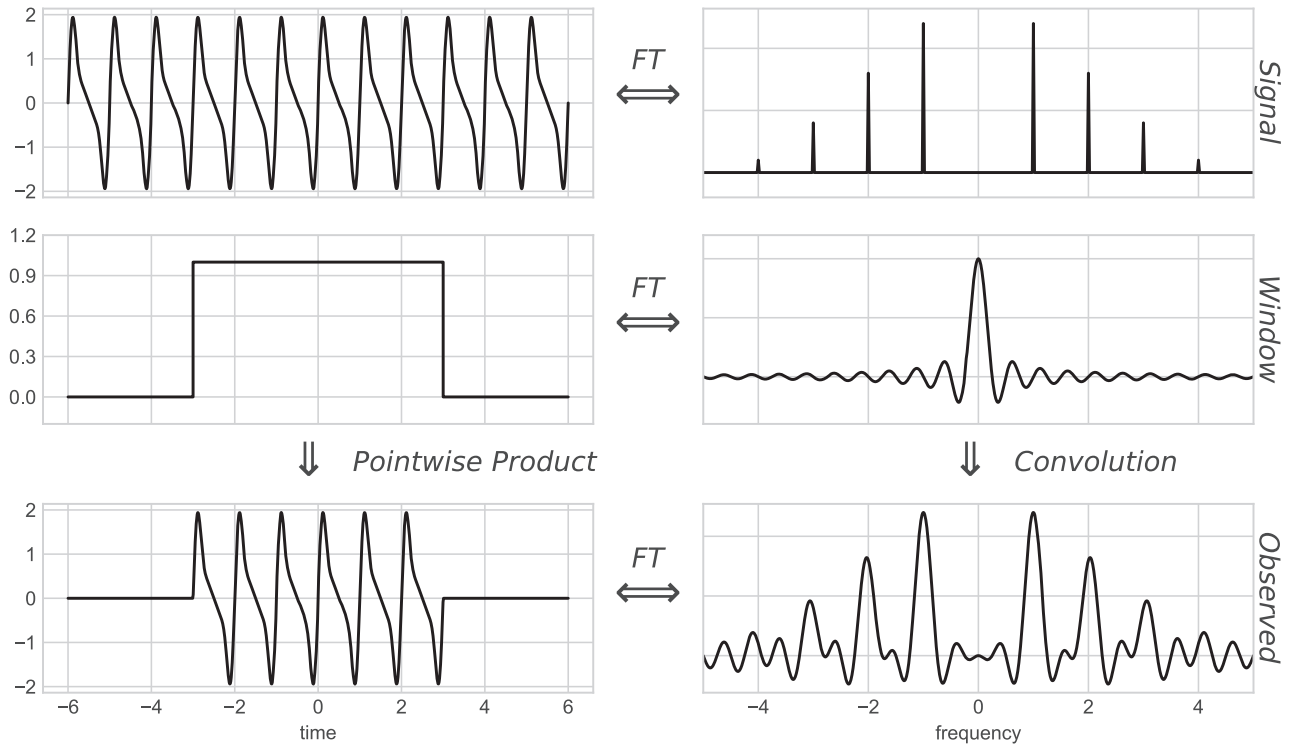


Figure 6. Visualization of the effect on the Fourier transform of a rectangular observing window (i.e., a continuous signal observed in its entirety within a finite range of time). The function used here is $g(t) = 1.2\sin(2\pi t) + 0.8\sin(4\pi t) + 0.4\sin(6\pi t) + 0.1\sin(8\pi t)$. The observed Fourier transform is a convolution of the true transform (here a series of Delta functions indicating the component frequencies) and the window transform (here a narrow sinc function).

We will see that these properties of the Fourier transform become essential when thinking about frequency components of time-domain measurements.

3. Window Functions: From Idealized to Real-world Signals

Until now we have been discussing Fourier transforms of continuous signals that are well defined for all times $-\infty < t < \infty$. Real-world temporal measurements of a signal, however, only involve some finite span of time, at some finite rate of sampling. In either case, the resulting data can be described by a pointwise product of the true underlying continuous signal with a window function describing the observation. For example, a continuous signal measured over a finite duration is described by a rectangular window function spanning the duration of the observation, and a signal measured at regular intervals is described by a Dirac comb window function marking those measurement times.

The Fourier transform of measured data in these cases, then, is not the transform of the continuous underlying function, but rather the transform of the *pointwise product of the signal and the observing window*. Symbolically, if the signal is $g(t)$ and the window is $W(t)$, the observed function is

$$g_{\text{obs}}(t) = g(t)W(t), \quad (21)$$

and by the convolution theorem, its transform is a convolution of the signal transform and the window transform:

$$\mathcal{F}\{g_{\text{obs}}\} = \mathcal{F}\{g\} * \mathcal{F}\{W\}. \quad (22)$$

This has some interesting consequences for the use and interpretation of periodograms, as we shall see.

3.1. Effect of a Rectangular Window

First, let us consider the case of observing a continuous periodic signal over a limited span of time: Figure 6 shows a continuous periodic function observed only within the window $-3 < t < 3$. The observed signal in this case can be understood as the pointwise product of the underlying infinite periodic signal with a rectangular window function. By the convolution theorem, the Fourier transform will be given by the convolution of the transform of the underlying function (here a set of delta functions at the component frequencies) and the transform of the window function (here a sinc function). For a purely periodic signal like the one seen in Figure 6, this convolution has the effect of replacing each delta function with a sinc function. Because of the inverse relationship between the width of the window and the width of its transform (see Figure 3), it follows that a wider observing window leads to proportionally less spread in each peak of the observed Fourier transform.

3.2. The Dirac Comb and the Discrete Fourier Transform

Another window function that commonly arises is when a continuous signal is sampled (nearly) instantaneously at regular intervals. Such an observation can be thought of as a pointwise product between the true underlying signal and a Dirac comb with the T parameter matching the spacing of the observations; this is illustrated in Figure 7. Interestingly, because the Fourier transform of a Dirac comb is another Dirac comb, the effect of such an observing window is to create a long sequence of




Figure 7. Visualization of the effect on the Fourier transform of a Dirac Comb observing window (i.e., a long string of evenly spaced discrete observations). The observed Fourier transform is a convolution of the true transform (here a localized Gaussian) and the window transform (here another Dirac comb).

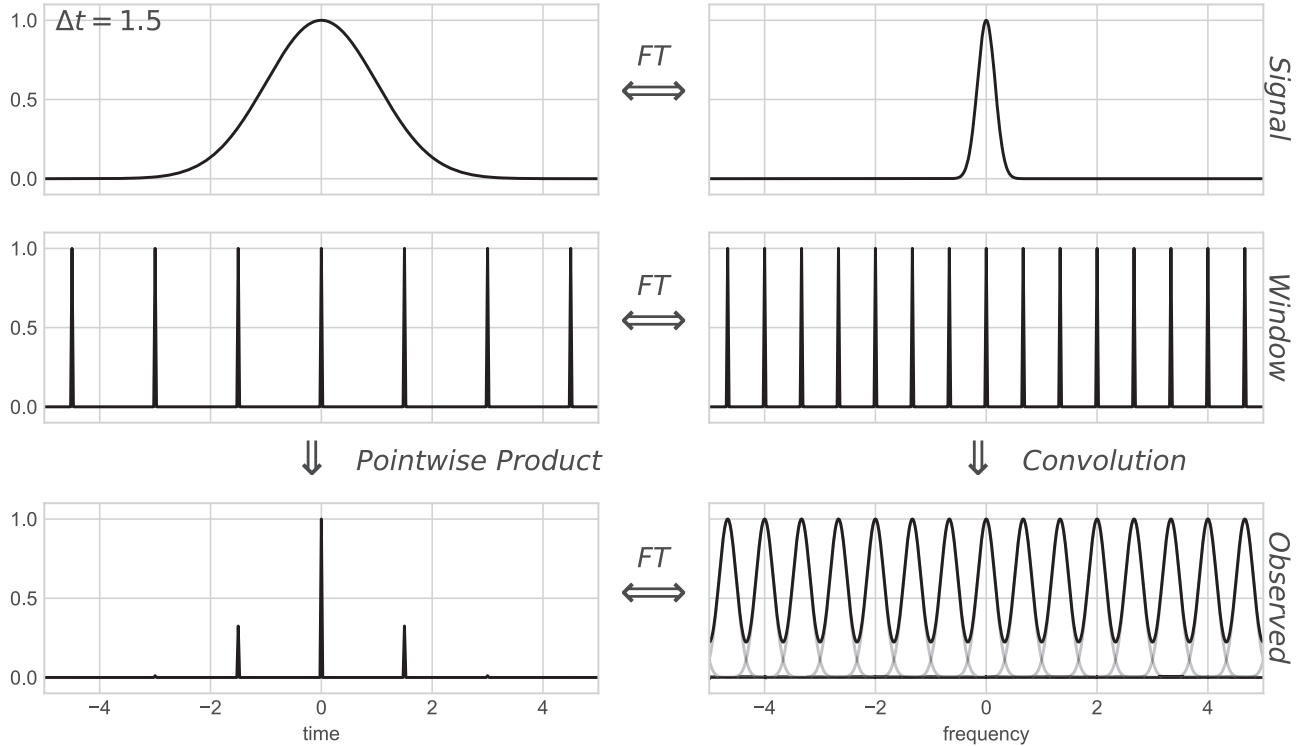


Figure 8. Repeating the visualization from Figure 7, but here with a lower sampling rate. The result is that the Fourier transform of the window function (middle right) has spacing narrower than the Fourier transform of the signal (top right), meaning that the observed Fourier transform (bottom right) has aliasing of signals, such that not all frequency information can be recovered. This is the reason for the famous Nyquist sampling theorem, which conceptually says that only a function whose Fourier transform can fit entirely between the “teeth” of the comb is able to be fully recovered for regularly spaced observations.

aliases of the underlying transform with a spacing of $1/T$. With this in mind, we can be assured in this case that evaluating the observed transform in the range $0 \leq f < 1/T$ is sufficient to

capture all the available frequency information: the signal outside that range is a sequence of identical aliases of what lies within that range.

3.2.1. The Nyquist Limit

The example in Figure 7 is somewhat of a best-case scenario, because the true Fourier transform values are nonzero only within a range of width $1/T$. If we increase the time between observations, decreasing the spacing of the frequency comb, the true transform no longer “fits” inside the window transform, and we will have a situation similar to that in Figure 8. The result is a mixing of different portions of the signal, such that the true Fourier transform *cannot be recovered* from the transform of the observed data!

This implies that if we have a regularly sampled function with a sampling rate of $f_0 = 1/T$, we can only fully recover the frequency information if the signal is *band limited* between frequencies $\pm f_0/2$. This is one way to motivate the famous Nyquist sampling limit, which approaches the question from the other direction and states that to fully represent the frequency content of a “band-limited signal” whose Fourier transform is zero outside the range $\pm B$, we must sample the data with a rate of at least $f_{\text{Ny}} = 2B$.

3.2.2. The Discrete Fourier Transform

When a continuous function is sampled at regular intervals, the delta functions in the Dirac comb window serve to collapse the Fourier integral into a Fourier sum, and in this manner we can arrive at the common form of the discrete Fourier transform. Suppose we have a true (infinitely long and continuous) signal $g(t)$, but we observe it only at a regular grid with spacing Δt . In this case, our observed signal is $g_{\text{obs}} = g(t)\text{III}_{\Delta t}(t)$ and its Fourier transform is

$$\hat{g}_{\text{obs}}(f) = \sum_{n=-\infty}^{\infty} g(n\Delta t)e^{-2\pi ifn\Delta t}, \quad (23)$$

which follows directly from Equations (1) and (17).

In the real world, however, we will not have an infinite number of observations, but rather a finite number of samples N . We can choose the coordinate system appropriately and define $g_n \equiv g(n\Delta t)$ to write

$$\hat{g}_{\text{obs}}(f) = \sum_{n=0}^N g_n e^{-2\pi ifn\Delta t}. \quad (24)$$

From the arguments around Nyquist aliasing, we know that the only relevant frequency range is $0 \leq f \leq 1/\Delta t$, and so we can define N evenly spaced frequencies with $\Delta f = 1/(N\Delta t)$ covering this range. Denoting the sampled transform as $\hat{g}_k \equiv \hat{g}_{\text{obs}}(k\Delta f)$, we can write

$$\hat{g}_k = \sum_{n=0}^N g_n e^{-2\pi i kn/N}, \quad (25)$$

which you might recognize as the standard form of the discrete Fourier transform.

Notice, though, that we glossed over one important thing: the effect of switching from an infinite number of samples to a finite number of samples. In moving from Equation (23) to Equation (24), we have effectively applied to our data a rectangular window function of width $N\Delta t$. From the discussion accompanying Figure 6, we know the result: it yields a Fourier transform convolved with a sinc function of width $1/(N\Delta t)$, resulting in the “smearing” of the Fourier transform signal with this width. Roughly speaking, then, any two Fourier transform values at frequencies within $1/(N\Delta t)$ of

each other will not be independent, and so we should space our evaluations of the frequency with $\Delta f \geq 1/(N\Delta t)$. Comparing to above, we see that this is *exactly the frequency spacing* we arrived at from Nyquist-frequency arguments.

What this indicates is that the frequency spacing of the discrete Fourier transform is optimal in terms of both the Nyquist sampling limit and the effect of the finite observing window! Now, this argument has admittedly been a bit hand-wavy, but there do exist mathematically rigorous approaches to proving that the discrete Fourier transform in Equation (25) captures all of the available frequency information for a uniformly sampled function g_n (see, e.g., Vetterli et al. 2014). Despite our lack of rigor here, I find this to be a helpful approach in developing intuition regarding the relationship between the continuous and discrete Fourier transforms.

3.3. The Classical Periodogram

With the discrete Fourier transform defined in Equations (24)–(25), we can apply the definition of the Fourier power spectrum from Equation (9) to compute the *classical periodogram*, sometimes called the *Schuster periodogram* after Schuster (1898), who first proposed it:

$$P_S(f) = \frac{1}{N} \left| \sum_{n=1}^N g_n e^{-2\pi ift_n} \right|^2. \quad (26)$$

Apart from the $1/N$ proportionality, this sum is precisely the Fourier power spectrum in Equation (9), computed for a continuous signal observed with uniform sampling defined by a Dirac comb. It follows that, in the uniform sampling case, the Schuster periodogram captures all of the relevant frequency information present in the data. This definition readily generalizes to the nonuniform case, which we will explore in the following section.

One point that should be emphasized is that the *periodogram* in Equation (26) and the *power spectrum* in Equation (9) are conceptually different things. As noted in Scargle (1982), the astronomy community tends to use these terms interchangeably, but to be precise, the periodogram (i.e., the statistic we compute from our data) is an *estimator* of the power spectrum (i.e., the underlying continuous function of interest). In fact, the classical periodogram and its extensions (including the Lomb–Scargle periodogram that we will discuss momentarily) are not consistent estimators of the power spectrum—that is, the periodogram has unavoidable intrinsic variance, even in the limit of an infinite number of observations (for a detailed discussion, see Chapter 8.4 of Anderson 1971).

4. Nonuniform Sampling

In the real world, particularly in fields like astronomy where observations are subject to influences of weather and diurnal, lunar, or seasonal cycles, the sampling rate is generally far from uniform. Using the same approach as we used to explore uniform sampling in the previous section, we can now explore nonuniform sampling here.

In the general nonuniform case, we measure some signal at a set of N times that we will denote $\{t_n\}$, which leads to the following observing window:

$$W_{\{t_n\}}(t) = \sum_{n=1}^N \delta(t - t_n). \quad (27)$$




Figure 9. Effect of nonuniform sampling on the observed Fourier transform. These samples have the same average spacing as those in Figure 7, but the irregular spacing within the observing window translates to irregular frequency peaks in its transform, causing the observed transform to be “noisy.” Here black and gray lines represent the real and imaginary parts of the transform, respectively.

Applying this window to our true underlying signal $g(t)$, we find an observed signal of the form

$$\begin{aligned} g_{\text{obs}}(t) &= g(t) W_{\{t_n\}}(t) \\ &= \sum_{n=1}^N g(t_n) \delta(t - t_n). \end{aligned} \quad (28)$$

Just as in the evenly sampled case, the Fourier transform of the observed signal is a convolution of the transforms of the true signal and the window:

$$\mathcal{F}\{g_{\text{obs}}\} = \mathcal{F}\{g\} * \mathcal{F}\{W_{\{t_n\}}\}. \quad (29)$$

Unlike in the uniform case, the window transform $\mathcal{F}\{W_{\{t_n\}}\}$ will generally *not* be a straightforward sequence of delta functions; the symmetry present in the Dirac comb is broken by the uneven sampling, leading the transform to be much more “noisy.” This can be seen in Figure 9, which shows the Fourier transform of a nonuniform observing window with an average sampling rate identical to that in Figure 7, along with its impact on the observed Fourier transform.

A few things stand out in this figure. In particular, the Fourier transform of the nonuniformly spaced delta functions looks like random noise, and in some sense it is: the locations and heights of the Fourier peaks are related to the intervals between observations, and so randomization of observation times leads to a randomization of Fourier peak locations and heights. In other words, *nonstructured spacing of observations leads directly to nonstructured frequency peaks in the window transform*. This nonstructured window transform, when convolved with the Fourier transform of the true signal, results in an observed Fourier transform reflecting the same random

noise. Comparing to the uniformly spaced observations in Figure 7, we see that the unstructured nature of the window transform means that there is no exact aliasing of the true signal and thus no way to exactly recover any portion of the true Fourier transform for the underlying function.

One might hope that sampling the signal more densely might alleviate these problems, and it does, but only to a degree. In Figure 10 we increase the density of observations by a factor of 10, such that there are 200 total observations over the length-10 observing window. The observed Fourier transform in this case is much more reflective of the underlying signal, but it still contains a degree of “noise” rooted in the randomized frequency peaks owing to randomized spacing between observations.

4.1. A Nonuniform Nyquist Limit?

We saw in Section 3.2.1 that the Nyquist limit is a direct consequence of the symmetry in the Dirac comb window function that describes evenly sampled data, and uneven sampling destroys the symmetry that underlies its definition. Nevertheless, the idea of the “Nyquist frequency” seems to have taken hold in the scientific psyche to the extent that the idea is often misapplied in areas where it is mathematically irrelevant. For unevenly sampled data, the truth is that the “Nyquist limit” might or might not exist, and even in cases where it does exist, it tends to be far larger (and thus far less relevant) than in the evenly sampled case.

4.1.1. Incorrect Limits in the Literature

In the scientific literature it is quite common to come across various proposals for a Nyquist-like limit applied in the case of

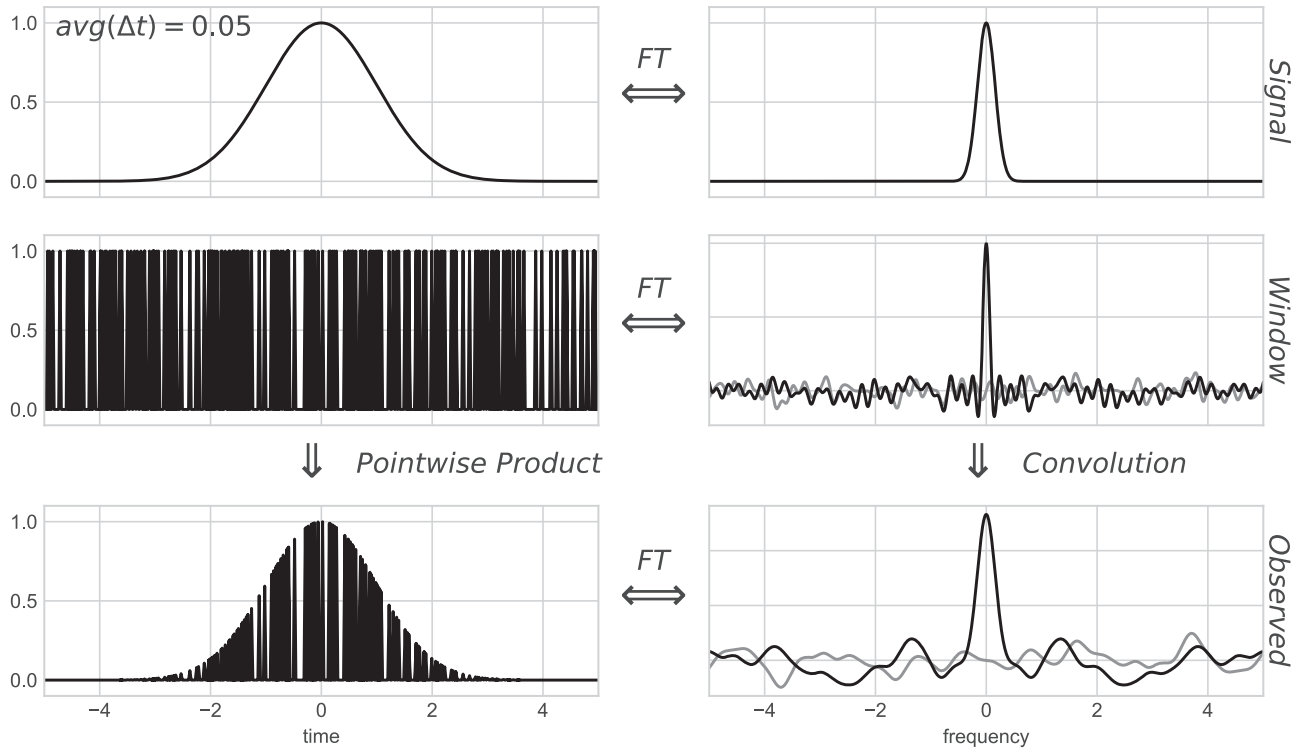


Figure 10. Effect of nonuniform sampling on the observed Fourier transform, with a factor of 10 more samples than Figure 9. Even with very dense sampling of the function, the Fourier transform cannot be exactly recovered owing to the imperfect aliasing present in the window transform.

irregular sampling. A few typical approaches include using the mean of the sampling intervals (e.g., Scargle 1982; Horne & Baliunas 1986; Press et al. 2007), the harmonic mean of the sampling intervals (e.g., Debosscher et al. 2007), the median of the sampling intervals (e.g., Graham et al. 2013a), or the minimum sample spacing (e.g., Percy 1986; Roberts et al. 1987; Press & Rybicki 1989; Hilditch 2001). All of these “pseudo-Nyquist” limits are tempting criteria in that they are easy to compute and reduce to the classical Nyquist frequency in the limit of evenly spaced data. Unfortunately, none of these approaches are correct: in general, unevenly sampled data can probe frequencies far larger than any of these supposed limits (a fact that several of these citations do hint at parenthetically).

As a simple example of where such logic can fail spectacularly, consider the data from Figure 1: though the mean sample spacing is one observation every 7 days, in Figure 2 we were nevertheless able to quite clearly identify a period of 2.58 hr—an order of magnitude shorter than the average-based pseudo-Nyquist limit would indicate as possible. For the data in Figure 1, the minimum sample spacing is just under 10 s, but it would be irresponsible to claim that this single pair of observations by itself defines some limit beyond which frequency information is unattainable.

As a more extreme example, consider the data shown in Figure 11. This consists of noisy samples from a sinusoid with a period of 0.01 units, with sample spacings ranging between 2 and 18 units: needless to say, any pseudo-Nyquist definition based on an average or minimum sample spacings will be *far* below the true frequency of 100; still, the the Lomb–Scargle periodogram in the bottom right panel quite cleanly recovers the true frequency.

4.1.2. The Nonuniform Nyquist Limit

While pseudo-Nyquist arguments based on average or minimum sampling fail spectacularly, there is a sense in which the Nyquist limit can be applied to unevenly spaced data. Eyer & Bartholdi (1999) explore this issue in some detail, and in particular they prove the following:

Let p be the largest value such that each t_i can be written $t_i = t_0 + n_i p$, for integers n_i . The Nyquist frequency then is $f_{\text{Ny}} = 1/(2p)$.

In other words, computing the Nyquist limit for unevenly spaced data requires finding the largest factor p , such that each spacing Δt_i is *exactly* an integer multiple of this factor. Eyer & Bartholdi (1999) prove this formally, but the result can be understood by thinking of such data as a windowed version of uniformly sampled data with spacing p , where the window is zero at all points except the location of the observations. Such uniform data have a classical Nyquist limit of $1/(2p)$, and a window function applied on top of that sampling does not change that fact.

Figure 12 shows an example of such a Nyquist frequency. The data are nonuniformly sampled at times $t_i = n_i \cdot p$, with $p = 0.01$ and n_i drawn randomly from positive integers less than 10,000. According to the Eyer & Bartholdi (1999) definition, this results in a Nyquist frequency $f_{\text{Ny}} = 50$, and we see the expected behavior beyond this frequency: the signal at $f > f_{\text{Ny}}$ consists of a series of exact aliases of the signal at $|f| < f_{\text{Ny}}$.

We should keep in mind one consequence of this Nyquist definition: if you have any pair of observation spacings whose ratio is irrational, *the Nyquist limit does not exist!* To realize this situation in practice, however, would require infinitely




Figure 11. Example of data for which the various poorly motivated “pseudo-Nyquist” approaches outlined in Section 4.1 fail spectacularly. The top panels show the data, a noisy sinusoid with a frequency of 100 (i.e., a period of 0.01). The bottom left panel shows a histogram of spacings between observations: the minimum spacing is 2.55, meaning that the signal has *over 250 full cycles* between the closest pair of observations. Nevertheless, the periodogram (bottom right) clearly identifies the correct period, though it is orders of magnitude larger than pseudo-Nyquist estimates based on an average or minimum sampling rate.

precise measurements of the times t_i ; finite precision of time measurements means that the Nyquist frequency can be large, but not infinite. For example, if your observation times are recorded to D decimal places, the Nyquist frequency will be at most

$$f_{\text{Ny}} \leq \frac{1}{2} 10^D, \quad (30)$$

with the inequality being due to the fact that larger common factors may exist. In other words, absent other relevant patterns in the observations, the Nyquist frequency for irregularly sampled data is most typically set by *the precision of the time measurements* (see also Bretthorst 2003; Koen 2006, for more rigorous treatments of this result).

4.1.3. Frequency Limit due to Windowing

In contexts where observations are not instantaneous, but rather consist of short-duration integrations of a continuous signal, a qualitatively different kind of frequency limit exists. This is typical in, e.g., optical astronomy, where a single observation typically consists of an integration of observed photons over a finite duration δt . As noted by Ivezić et al. (2014), this timescale of integration represents another kind of limiting frequency for irregularly sampled data. Such a situation means that the observation is effectively a convolution of the underlying signal with a rectangular window function of width δt , in a manner analogous to Figure 5. By the

convolution theorem, the observed transform will be a pointwise product between the true transform and the transform of the window, which will generally have a width proportional to $1/\delta t$. This means that—absent other more constraining window effects—the frequency limit is $f_{\text{max}} \propto 1/(2\delta t)$, with the constant of proportionality dependent on the shape of the effective window describing individual observations.

Keep in mind that the windowing limit $1/(2\delta t)$ is quite different from a Nyquist limit: the Nyquist limit is the frequency beyond which all signal is aliased into the Nyquist range; the windowing limit is the frequency beyond which all signal is attenuated to zero. In practice, the limit implied by either the temporal resolution or windowing of individual observations may be too large to be computationally feasible; for discussion of frequency limits in practice, see Section 7.1.

4.2. Semistructured Observing Windows

We have seen that for uniform data the perfect aliasing beyond the Nyquist frequency is a direct consequence of the symmetry of the Dirac-comb window function. For nonuniform observations, such symmetry does not exist, but *structure* in the observing window can lead to partial aliasing of signals in the data (see, e.g., Deeming 1975). In this section, we will examine two typical window functions derived from real-world observations: one ground-based (LINEAR) and one space-based (*Kepler*). For details on how window power spectra can be estimated in practice, see Section 7.3.

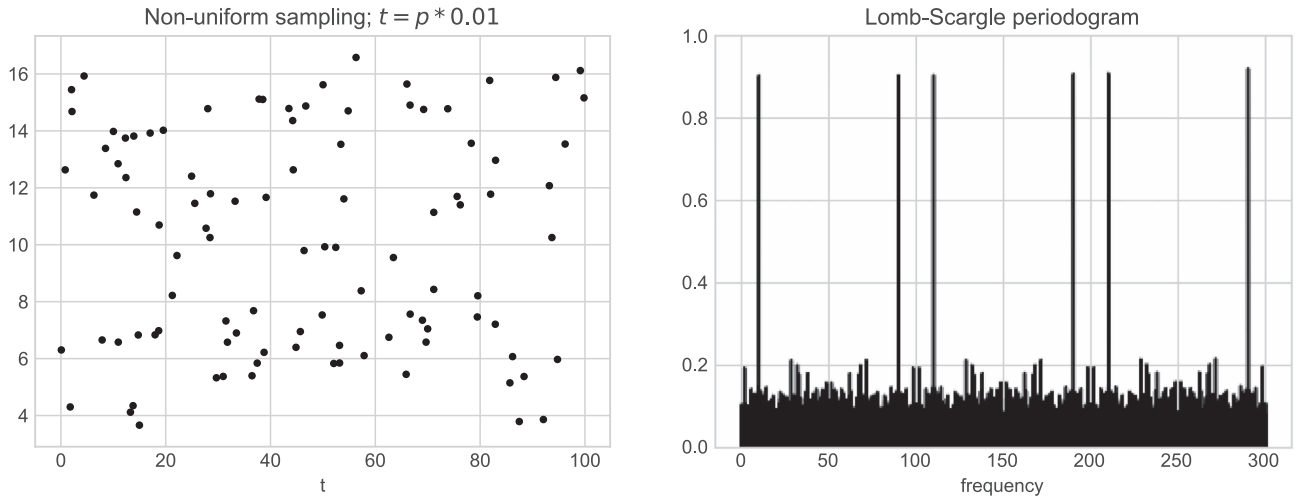


Figure 12. Visualization of the Eyer & Bartholdi (1999) definition of the Nyquist frequency. Data are nonuniformly sampled at times $t_i = n_i p$, for integer n_i and $p = 0.01$. This results in a Nyquist frequency of $f_{\text{Ny}} = (2p)^{-1} = 50$: the periodogram outside the range $0 \leq f < f_{\text{Ny}}$ is a series of perfect aliases of the signal within that range.




Figure 13. Power spectrum of the observing window for the data shown in Figure 1. Notice the strong spike in power at a period of 1 day and related aliases at $1/n$ days for integer n . There is also a strong spike at 365 days and a noticeable spike at ~ 14 days. Each of these indicates time intervals that appear often in the data.

4.2.1. A Ground-based Observing Window: LINEAR

Let us again consider the data shown in Figure 1. The window power spectrum for this in Figure 13 shows some quite distinct features, and these features have an intuitive interpretation. Namely, if the window power shows a spike at a period of p days, this means that an observation at time t_0 is likely to be followed by another observation near a time $t_0 + np$ for integer n .

With this in mind, the strong spike at a period of 1 day indicates that observations are taken near the same time of day: this is typical of a ground-based survey with observations recorded only during the nighttime hours. The additional spikes at periods of $1/n$ days (for integer n) are aliases of this same feature. Figure 13 also shows a wide spike at a period of 1 yr, indicating a detectable annual pattern in the observations. Finally, there is a noticeable spike at approximately 14 days that is likely related to patterns of scheduling within the survey.

Recall that a Fourier spectrum observed through a particular window will reflect a convolution of the true spectrum and the window spectrum (see Figure 9), and so we would expect the structure in the window to be imprinted on the power spectrum measured from the data.

This imprint of the window power is illustrated in Figure 14. The top panel shows the window power spectrum as a function of frequency (rather than period, as in Figure 13), while the

bottom panel shows the observed signal power spectrum as a function of frequency (rather than period, as in Figure 2). The top panel and its inset show clearly the 1 day and 1 yr features we noted previously. The bottom panel shows the observed power spectrum of the data: these diurnal and annual peaks in the window function are quite clearly imprinted on the observed power spectrum at relevant scales. This approximate aliasing is similar to the exact aliasing seen in regularly sampled data at the Nyquist frequency; however, in this case the magnitude of the aliased signal fades further from the frequency driving the signal.

4.2.2. A Space-based Observing Window: Kepler

Space-based surveys will generally have quite different observing windows. For example, the top left panel of Figure 15 shows observations of an RR Lyrae variable star from the *Kepler* survey, measured 4083 times over a period of 3 months, with an irregular observing cadence of around 30 minutes (for a deeper discussion of these observations, see Kolenberg et al. 2010). The *Kepler* observations are very nearly uniformly spaced, and this is reflected in the window power spectrum, shown in the top right panel of Figure 15. The window function is a series of very narrow evenly spaced spikes, reminiscent of the Dirac comb shown in Figures 7–8. By analogy we can treat $f_{\text{Ny}} = 0.5/29.4 \text{ minute}^{-1}$ as the

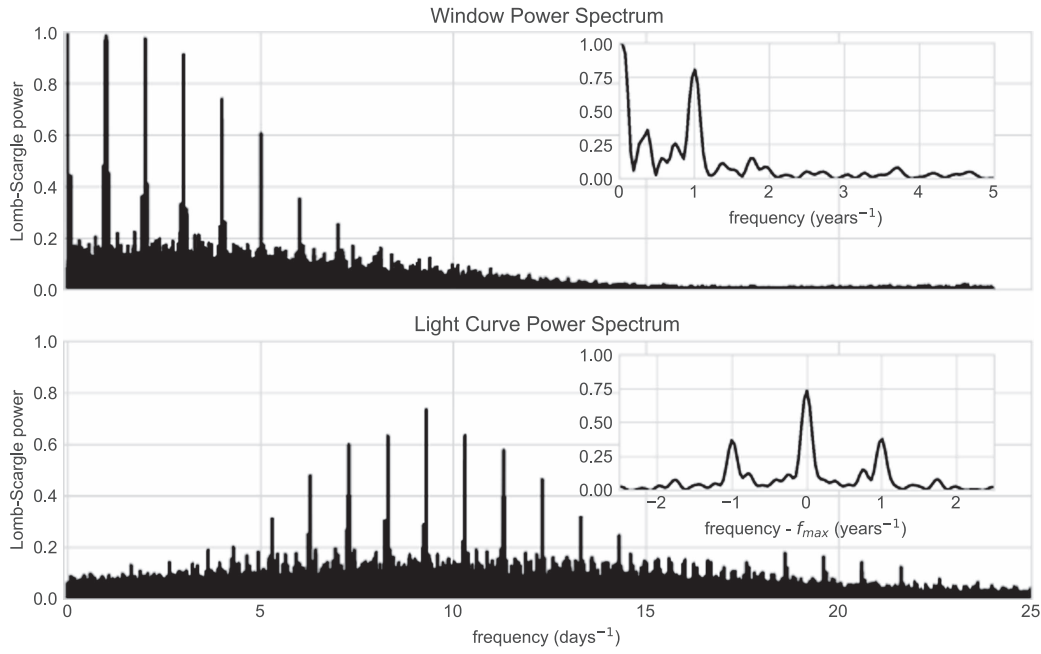


Figure 14. Effect of the window function in Figure 13 on the power spectrum in Figure 2. The top panel shows the window power spectrum, and the bottom panel shows the observed signal power spectrum. Both are plotted as a function of frequency (we earlier saw both of these as a function of period; see Figures 13 and 2, respectively). Viewing these as a function of frequency makes it clear that the structure in the window function is imprinted on the observed spectrum: both the diurnal structure in the main panel and the annual structure in the inset are apparent in the observed spectrum.

effective Nyquist limit for the data, keeping in mind that aliasing beyond this “limit” will be imperfect owing to the uneven spacing of the samples (see the bottom left panels of Figure 15). The bottom right panel of Figure 15 shows the power spectrum of the observations, with gray shading indicating the (nearly) aliased region of the spectrum. The period of 13.6 hr is quite strongly apparent, along with smaller spikes at integer multiples of this frequency that indicate higher-order periodic components in the signal.

The window functions for ground-based and space-based observations, reflected by LINEAR data in Figure 14 and *Kepler* data in Figure 15, are quite different, but in both cases essential features of the observed power spectra can be understood by recognizing that the periodogram measures not the power spectrum of the underlying signal but a power spectrum from the convolution of the true signal transform and the Fourier transform of the window function.

5. From Classical to Lomb–Scargle Periodograms

Up until now, we have been mainly discussing the direct extension of the classical periodogram in Equation (26) to nonuniform data. Returning to this definition, we can rewrite the expression in a more suggestive way:

$$\begin{aligned}
 P(f) &= \frac{1}{N} \left| \sum_{n=1}^N g_n e^{-2\pi i f t_n} \right|^2 \\
 &= \frac{1}{N} \left[\left(\sum_n g_n \cos(2\pi f t_n) \right)^2 + \left(\sum_n g_n \sin(2\pi f t_n) \right)^2 \right].
 \end{aligned} \tag{31}$$

Although this form of the nonuniform periodogram can be useful for identifying periodic signals, its statistical properties are not as straightforward as in the uniform case. When the classical periodogram is applied to uniformly sampled Gaussian white noise, the values of the resulting periodogram are chi-square distributed. This property becomes quite useful in practice when the periodogram is used in the context of a classical hypothesis test to distinguish between periodic and nonperiodic objects—see Section 7.4.2. Unfortunately, when the sampling becomes nonuniform, these properties no longer hold and the periodogram distribution cannot in general be analytically expressed.

Scargle (1982) addressed this by considering a generalized form of the periodogram,

$$\begin{aligned}
 P(f) &= \frac{A^2}{2} \left(\sum_n g_n \cos(2\pi f [t_n - \tau]) \right)^2 \\
 &\quad + \frac{B^2}{2} \left(\sum_n g_n \sin(2\pi f [t_n - \tau]) \right)^2,
 \end{aligned} \tag{32}$$

where A , B , and τ are arbitrary functions of the frequency f and observing times $\{t_i\}$ (but not the values $\{g_n\}$), and showed that you can choose a unique form of A , B , and τ such that

1. the periodogram reduces to the classical form in the case of equally spaced observations;
2. the periodogram’s distribution is analytically computable; and
3. the periodogram is insensitive to global time shifts in the data.

Kepler object ID 007198959




Figure 15. RR Lyrae variable observed by the *Kepler* project (see Kolenberg et al. 2010). Top left: 4083 observed fluxes over roughly 3 months. Top right: window power spectrum, which is quite close to that of regularly spaced data with a cadence of 29.4 minutes. Bottom left: time between observations. Missing measurements aside, the spacing between observations is nearly uniform. The lower panel gives a closer look at the majority of the spacings, which are not exactly the same but rather span a range of ± 50 ms around the 29.4-minute period observed in the window function. Bottom right: data power spectrum, showing approximate aliasing and clearly indicating a peak near a period of 13.6 hr, along with higher-order components at multiples of this value.




Figure 16. Top panel: comparison of the classical periodogram (Equation (31)) and the Lomb–Scargle periodogram (Equation (33)) for 30 noisy points drawn from a sinusoid. Though the two periodogram estimates differ quantitatively, the essential qualitative features—namely, the position of significant peaks—typically remain the same. Bottom panel: values of the denominators in Equation (33). The difference between the Lomb–Scargle periodogram and the classical periodogram stems from the difference between these quantities and $N/2 = 15$ (dotted line).

The values of A and B leading to these properties result in the following form of the generalized periodogram:

$$P_{LS}(f) = \frac{1}{2} \left\{ \left(\sum_n g_n \cos(2\pi f[t_n - \tau]) \right)^2 / \sum_n \cos^2(2\pi f[t_n - \tau]) + \left(\sum_n g_n \sin(2\pi f[t_n - \tau]) \right)^2 / \sum_n \sin^2(2\pi f[t_n - \tau]) \right\}, \quad (33)$$

where τ is specified for each f to ensure time-shift invariance:

$$\tau = \frac{1}{4\pi f} \tan^{-1} \left(\frac{\sum_n \sin(4\pi f t_n)}{\sum_n \cos(4\pi f t_n)} \right). \quad (34)$$

This modified periodogram differs from the classical periodogram only to the extent that the denominators $\sum_n \sin^2(2\pi f t_n)$ and $\sum_n \cos^2(2\pi f t_n)$ differ from $N/2$, which is the expected value of each of these quantities in the limit of complete phase sampling at each frequency. Thus, in many cases of interest the Lomb–Scargle periodogram only differs slightly from the classical/Schuster periodogram; an example of this is seen in Figure 16.

A remarkable feature of Scargle’s modified periodogram is that it is *identical* to the result obtained by fitting a model consisting of a simple sinusoid to the data at each frequency f and constructing a “periodogram” from the χ^2 goodness of fit at each frequency—an estimator that was considered in some depth by Lomb (1976). From this perspective, the τ shift defined in Equation (34) serves to orthogonalize the normal




Figure 17. Sinusoidal model implied by the Lomb–Scargle periodogram for the LINEAR data seen in Figure 1. Although a sinusoid does not perfectly fit the data, the sinusoidal model is close enough that it serves to locate the correct frequency.

equations used in the least-squares analysis. Partly due to this deep connection between Fourier analysis and least-squares analysis, the modified periodogram in Equation (33) has since become commonly referred to as the *Lomb–Scargle periodogram*, although versions of this approach had been employed even earlier (see, e.g., Gottlieb et al. 1975).

Because of the close similarity between the classical and Lomb–Scargle periodograms, the bulk of our previous discussion applies—at least qualitatively—to periodograms computed with the Lomb–Scargle method. In particular, reasoning about the effect of window functions on the observed Lomb–Scargle power spectrum remains qualitatively useful even if it is not quantitatively precise.

One important caveat of the simple Lomb–Scargle formula is that the statistical guarantees only apply when the observations have *uncorrelated white noise*; data with more complicated noise characteristics must be treated more carefully (see, e.g., Vio et al. 2010, or the least-squares approach discussed in Section 6.1).

6. The Least-squares Periodogram and Its Extensions

The equivalence of the Fourier interpretation and least-squares interpretation of the Lomb–Scargle periodogram allows for some interesting and useful extensions, some of which we will explore in this section. First, let us consider the least-squares periodogram itself.

In the least-squares interpretation of the periodogram, a sinusoidal model is proposed at each candidate frequency f :

$$y(t; f) = A_f \sin(2\pi f(t - \phi_f)), \quad (35)$$

where the amplitude A_f and phase ϕ_f can vary as a function of frequency. These model parameters are fit to the data in the standard least-squares sense, by constructing the χ^2 statistic at each frequency:

$$\chi^2(f) \equiv \sum_n (y_n - y(t_n; f))^2. \quad (36)$$

We can find the “best” model $\hat{y}(t; f)$ by minimizing $\chi^2(f)$ at each frequency with respect to A_f and ϕ_f ; we will denote this minimum value as $\hat{\chi}_0^2(f)$. Scargle (1982) showed that with this setup the Lomb–Scargle periodogram from Equation (33) can be equivalently written:

$$P(f) = \frac{1}{2} [\hat{\chi}_0^2 - \hat{\chi}^2(f)], \quad (37)$$

where $\hat{\chi}_0^2$ is the nonvarying reference model. The key realization here is that the Lomb–Scargle periodogram

essentially *assumes a sinusoidal model* for the data; this is visualized in Figure 17 for the data we had seen in Figure 1. This immediately begs the question, can we compute a “periodogram” based on more general forms of the above model to more effectively fit the data?

6.1. Measurement Uncertainties

Perhaps the most important modification to the periodogram is to construct it such that it correctly handles uncertainties in the observed data. This can be done through the standard change to the χ^2 expression, i.e., if there are Gaussian errors σ_n on each observed y_n , we can rewrite Equation (36) in the following way:

$$\chi^2(f) \equiv \sum_n \left(\frac{y_n - y_{\text{model}}(t_n; f)}{\sigma_n} \right)^2. \quad (38)$$

The periodogram constructed from this χ^2 definition will more accurately reflect the spectral power of noisy observations. The effect of this modification on Equation (33) is the addition of a multiplicative weight $1/\sigma_n$ within each of the summations. Early versions of this sort of modification appeared in Gilliland & Baliunas (1987) and Irwin et al. (1989), Scargle (1989) derived this “weighted” form of the periodogram without direct reference to the least-squares model, and Zechmeister & Kürster (2009) showed that such a modification does not change the statistical properties of the resulting periodogram.

This generalization of $\chi^2(f)$ also suggests a convenient way to construct a periodogram in the presence of correlated observational noise. If we let Σ denote the $N \times N$ noise covariance matrix for N observations and construct the vectors

$$\mathbf{y} = [y_1, y_2, \dots, y_N]^T$$

$$\mathbf{y}_{\text{model}} = [y_{\text{model}}(t_1), y_{\text{model}}(t_1), \dots, y_{\text{model}}(t_N)]^T, \quad (39)$$

then the χ^2 expression for correlated errors can be written as

$$\chi^2(f) = (\mathbf{y} - \mathbf{y}_{\text{model}})^T \Sigma^{-1} (\mathbf{y} - \mathbf{y}_{\text{model}}), \quad (40)$$

which reduces to Equation (38) if noise is uncorrelated (i.e., if the off-diagonal terms of Σ are zero). This resulting periodogram is quite similar to the approach to correlated noise developed by Vio et al. (2010) from the Fourier perspective, and is in fact exactly equivalent in the case of the “zero-mean colored noise” example considered therein.

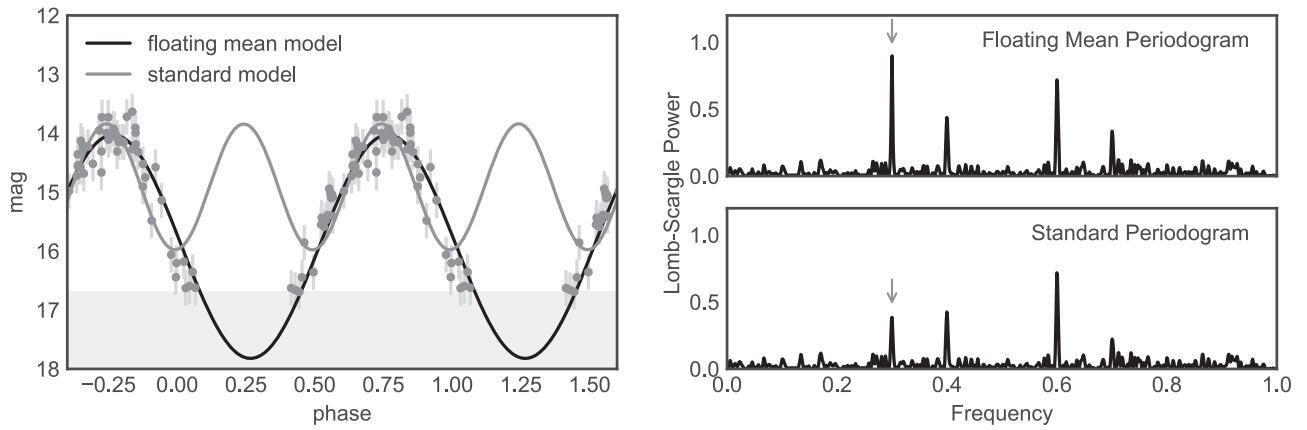


Figure 18. Comparison of the standard and floating-mean periodograms for data with a frequency of 0.3 and a selection effect that removes faint observations. In this case the mean estimated from the observed data is not close to the true mean, which leads to the failure of the standard periodogram to recover the correct frequency. A floating-mean model correctly recovers the true frequency of 0.3.

6.2. Data Centering and the Floating-mean Periodogram

Another commonly applied modification of the periodogram has variously been called the *date-compensated discrete Fourier transform* (Ferraz-Mello 1981), the *floating-mean periodogram* (Cumming et al. 1999; VanderPlas & Ivezić 2015), or the *generalized Lomb-Scargle method* (Zechmeister & Kürster 2009) and involves adding an offset term to the sinusoidal model at each frequency:³

$$y_{\text{model}}(t; f) = y_0(f) + A_f \sin(2\pi f(t - \phi_f)). \quad (41)$$

This turns out to be quite important in practice because the standard Lomb-Scargle approach assumes that the data are pre-centered around the mean value of the (unknown) signal. In many analyses, this requirement is accomplished by pre-centering data about the sample mean: this approach is generally robust when the data provide full phase coverage of the observed signal; however, due to selection effects and survey cadence, full phase coverage cannot always be guaranteed. Using the sample mean in such cases can potentially lead to suppression of peaks of interest.

A simulated example of this is shown in Figure 18: the data consist of noisy observations of a sinusoidal signal in which the faintest observations are omitted from the data set (due to, e.g., a detection threshold). Applying the standard Lomb-Scargle periodogram to pre-centered data leads to a periodogram that suppresses the true period of 0.3 days (bottom right panel). Using the floating-mean model of Equation (41) yields a periodogram that identifies this true period (top right panel). A detailed study of the floating-mean model is given by Zechmeister & Kürster (2009), who show that the addition of the floating-mean term does not change the useful statistical properties outlined in Section 5.

6.3. Higher-order Fourier Models

A further generalization of the least-squares periodogram involves multiterm Fourier models: rather than fitting just a single sinusoid at each frequency, we might fit a partial Fourier series, adding $K - 1$ additional sinusoidal components at

integer multiples of the fundamental frequency:

$$y_{\text{model}}(t; f) = A_f^0 + \sum_{k=1}^K A_f^{(k)} \sin(2\pi k f(t - \phi_f^{(k)})). \quad (42)$$

Bretthorst (1988) takes a comprehensive look at this type of multiterm extension to the periodogram, as well as related extensions to decaying signals and other more complex models.

This kind of Fourier series generalization to the periodogram is quite tempting, because it means that the periodogram can be tuned to fit models that are more complicated than simple sine waves. In some cases, this can be very useful; for example, Figure 19 shows a Lomb-Scargle analysis of an eclipsing binary star, characterized by both a primary and secondary eclipse. The standard Lomb-Scargle periodogram (top panels) is maximized at twice the true rotation frequency, because the simple sinusoidal model is unable to closely fit the primary and secondary eclipses separately. A six-term Fourier model (bottom panels) is sufficiently flexible that it can detect the true 17.5 hr period, though this comes at the expense of a much noisier periodogram.

This additional periodogram noise is easy to understand: in the least-squares view of the periodogram, the periodogram height at any frequency is directly related to how well the model fits the data. For nested linear models, adding additional terms will *always* provide a fit to data equal to or better than the simpler model, and so it follows that a periodogram based on a more complex model will be higher *at all frequencies*, not just at frequencies of interest! Indeed, this effect can be readily observed in Figure 19.

This background noise in the multiterm periodogram can also be understood in terms of aliasing. Consider Figure 20, which shows several multiterm fits to the *Kepler* data from Figure 15. Given a standard periodogram with a peak or subpeak at f_0 , a two-term periodogram will duplicate this peak at $f_0/2$, with the second harmonic driving the fit. Similarly, a three-term periodogram will add additional peaks at both $f_0/2$ and $f_0/3$, as a result of the original peak falling in the second and third harmonic. In general, you can expect an N -term periodogram to contain N aliases of every feature present in the standard periodogram; any strong peak revealed by the multiterm periodogram is due to two or more of these aliases coinciding.

³ We choose to follow Cumming et al. (1999) and VanderPlas & Ivezić (2015) and call this a “floating-mean” model, to avoid confusion of the different uses of the term “generalized periodogram” in, e.g., Bretthorst (2001) and Zechmeister & Kürster (2009).




Figure 19. One-term and six-term Lomb–Scargle models fit to LINEAR object 14752041, an eclipsing binary. Notice that the standard periodogram (top panels) finds an alias of the true 17.5 hr frequency, because a simple sinusoidal model cannot closely fit both the primary and secondary eclipse. A six-term Fourier model, on the other hand, does find the true period (bottom panels).

In cases where the single-term power spectrum is itself noisy and/or dominated by partial aliasing due to window effects, these additional aliases can quickly wash out any gains from the more complicated model. For example, consider the multiterm periodograms for the LINEAR data depicted in Figure 21: as previously discussed, the LINEAR periodogram contains a number of aliases of the true peak owing to the dominant 1 day signal in the window function. Adding terms to this model only compounds this problem, increasing the number of spurious peaks at the low-frequency end. For data with even moderate levels of noise, chance coincidences of these peaks can lead to spurious detections that dominate the true peak, particularly for models with many Fourier terms.

6.4. Additional Extensions

When the periodogram is viewed from the least-squares model-fitting perspective, there is no need to limit the analysis to sums of sinusoids. There have been some very interesting applications extending this type of analysis to more arbitrary models. For example, periodogram models can be extended to measure decaying signals, nonstationary signals, multifrequency signals, chirps, and other signal types (see, e.g., Jaynes 1987; Bretthorst 1988; Gregory 2001). In particular, Bretthorst (1988) demonstrates the effectiveness of such approaches in applications ranging from medical imaging to astronomy. The challenge of such extensions is the fact that you often need to use some prior knowledge of the system being observed to decide whether a more complicated model is indicated, as well as what form of model to apply. In practice, this comes up only when searching for very specific signals for which a more complicated model has some a priori physical motivation.

On the astronomy side, several examples of this flavor of extension exist. For example, the Supersmoother approach to detecting periodicity involves fitting a flexible nonparametric smoothing function to the data at each frequency (Reimann 1994): the flexibility of the model leads to fewer aliasing issues when compared to the more constrained sinusoidal model. Another approach is to use empirically derived templates as a functional fit at each frequency; this has been employed effectively in Sesar et al. (2010, 2013) and related work.

Finally, there has been some exploration of extensions to the Lomb–Scargle periodogram for use with multiband observations, using various forms of regularization to control model complexity (VanderPlas & Ivezić 2015; Long et al. 2016). With many of these least-squares and/or Bayesian extensions, computational complexity quickly becomes an issue, because fast $\mathcal{O}(N \log N)$ approaches that can be used for sinusoidal fits (see Section 7.6) are not available for more general functional forms, though there is some promising work in this area: see, for example, the *Fast Template Periodogram*⁴ (J. A. Hoffman et al. 2017, in preparation), which can quickly construct periodograms from Fourier approximations to templates.

6.5. A Note about “Bayesian” Periodograms

The least-squares view of the Lomb–Scargle periodogram creates a natural bridge, via maximum likelihood, to Bayesian periodic analysis. In fact, Jaynes (1987) showed that the standard form of the Lomb–Scargle periodogram can be derived directly from the axioms of Bayesian probability theory outlined in his comprehensive treatment of the subject (Jaynes & Bretthorst 2003).⁵ In the Bayesian view, the Lomb–Scargle periodogram is in fact the optimal statistic for detecting a stationary sinusoidal signal in the presence of Gaussian noise.

For the standard, simple-sinusoid model, the Bayesian periodogram is given by the posterior probability of frequency f given the data D and sinusoidal model M :

$$p(f|D, M) \propto e^{P_{\text{LS}}(f)}, \quad (43)$$

where $P_{\text{LS}}(f)$ is the Lomb–Scargle power from Equation (33). The effect of this exponentiation, as seen in Figure 22, is to suppress sidelobes and alias peaks in relation to the largest one of the spectrum.

The benefit of the Bayesian approach is that it allows more flexible models, more principled treatment of nuisance parameters related to noise in the data, and the ability to make use of prior information in a periodic analysis. It explicitly returns probability density as a function of frequency, and it is tempting to think of this as the “natural” way to interpret the periodogram.

⁴ <http://ascl.net/code/v/1559>

⁵ Read this. Really.




Figure 20. Multiterm models (left) and their corresponding periodograms (right) for the *Kepler* data shown in Figure 15.

The problem, however, is that the Bayesian periodogram does *not* compute the probability that the data are periodic with a given frequency, but rather that probability conditioned on a relatively strong assumption: that *the data are drawn from a sinusoidal model*. As such, the standard Bayesian periodogram is not a useful measure of generic periodicity! In fact, the Bayesian periodogram’s derivation under the assumption of a sinusoidal signal is perhaps the best argument against its use for unknown signals: the result of a Bayesian analysis is only ever as good as the assumptions that go into it, and for a general (not necessarily sinusoidal) signal, those assumptions are incorrect, and the periodogram should not be trusted.

Now, the standard Lomb–Scargle periodogram can also be viewed as derived from a sinusoidal fit and thus might appear subject to the same criticism. But unlike the Bayesian periodogram, the standard periodogram affords interpretation in light of Fourier analysis and window functions: here the incorrect sinusoidal model manifests itself in terms of frequency aliases, which can be understood through the analysis of window functions. In most cases of interest, such analysis turns out to be *vital* to the application and interpretation of the periodogram (see Section 7.2).

In other words, the Bayesian approach essentially goes “all-in” on the least-squares interpretation of the periodogram, exponentially suppressing the information that allows you to reason about the periodogram in light of its connection to Fourier analysis. In contrived cases with clean sinusoidal data and unstructured window functions, exponential attenuation of sidelobes and aliases may seem appealing (see, e.g., Mortier et al. 2015), but that appeal extends to the real world only in the most favorable of cases—i.e., high signal-to-noise ratio measurements of near-sinusoidal data with a very well behaved survey window. In short, you should be wary of placing too much trust in a Bayesian periodogram, unless you are certain of the precise form of the signal you are looking for.

This is not to say that all Bayesian approaches to periodic analysis are similarly flawed; there have been many interesting studies that go beyond the simple sinusoidal model and use more complex and/or flexible models. Some examples are models based on Fourier extensions with strong priors (e.g., Bretthorst 1988), instrument-dependent parametric models (e.g., Angus et al. 2016), flexible nonparametric functions (e.g., Gregory & Loredó 1992), Gaussian process models (e.g., Wang et al. 2012), and specially designed stochastic models

(e.g., Kelly et al. 2014). Though these are often more accurate and powerful than the Lomb–Scargle approach, they tend to be far more expensive computationally. Bayesian approaches based on Markov Chain Monte Carlo (MCMC) also tend to run into stability problems, particularly for multimodal or other complicated posterior distributions (see, e.g., the RR Lyrae discussion in Kelly et al. 2014).

7. Practical Considerations When Using Lomb–Scargle Periodograms

The previous sections have given a conceptual introduction to the Lomb–Scargle periodogram and its roots in both Fourier and least-squares analysis. We now come to the meat of the subject at hand: given this understanding of the Lomb–Scargle periodogram and related approaches, how should we use it most effectively in practice? The following sections will identify several of the important issues and questions that are not often addressed in the literature on the subject but are nevertheless vital to consider when using the periodogram in practice.

7.1. Choosing a Frequency Grid

The frequency grid used for a Lomb–Scargle analysis is an important choice that is too often glossed over, probably because the choice is so straightforward in the more familiar case of uniformly sampled data. For nonuniform data, it is not so simple, and there are two important considerations: the *frequency limits* and the *grid spacing*.

The frequency limit on the low end is relatively easy: for a set of observations spanning a length of time T , a signal with frequency $1/T$ will complete exactly one oscillation cycle, and so this is a suitable minimum frequency. Often, it is more convenient just to set this minimum frequency to zero, as it does not add much of a computational burden and is unlikely to add any significant spurious peak to the periodogram.

The high-frequency limit is more interesting and goes back to the discussion of Nyquist and/or limiting frequencies from Section 4.1: in order not to miss relevant information, it is important to compute the periodogram up to some well-motivated limiting frequency f_{max} . This could be a true Nyquist limit based on the Eyer & Bartholdi (1999) definition, a pseudo-Nyquist limit based on careful scrutiny of the window function (see Figure 15), a limiting frequency based on the




Figure 21. Multiterm models (left) and their corresponding periodograms (right) for the LINEAR data shown in Figure 1.

integration time of individual observations, or a limit based on prior knowledge of the kinds of signals you expect to detect.

With the frequency range decided, we next must determine how finely to sample the frequencies between the limits. This choice turns out to be quite important as well: too fine a grid can lead to unnecessarily long computation times that can add up quickly in the case of large surveys, while too coarse a grid risks entirely missing narrow peaks that fall between grid points. For example, Figure 23 shows the true, well-sampled periodogram (gray line), along with a periodogram computed for 10,000 equally spaced periods covering the same range (black line). Because the spacing of the 10,000-point grid is much larger than the width of the periodogram peaks, the analysis *entirely misses* the most important peaks in the periodogram!

This shows us that it is important to choose grid spacings smaller than the expected widths of the periodogram peaks. From our discussion of windowing in Section 3 (particularly Figure 6), we know that data observed through a rectangular window of length T will have sinc-shaped peaks of width $\sim 1/T$. To ensure that our grid sufficiently samples each peak, it is prudent to oversample by some factor—say, n_o samples per peak—and use a grid of size $\Delta f = 1/n_o T$. This pushes the total number of required periodogram evaluations to

$$N_{\text{eval}} = n_o T f_{\text{max}}. \quad (44)$$

So what is a good choice for n_o ? Values ranging from $n_o = 5$ (Schwarzenberg-Czerny 1996; VanderPlas & Ivezić 2015) to $n_o = 10$ (Debusscher et al. 2007; Richards et al. 2012) seem to be common in the literature; for periodograms computed in this paper, we use $n_o = 5$.

Depending on the characteristics of the data set, the size of the resulting frequency grid can vary greatly. For example, the *Kepler* data shown in Figure 15 have a pseudo-Nyquist frequency of 48.9 day^{-1} and an observing window of $T = 90$ days. To compute five samples per peak thus requires $N_{\text{eval}} \approx 22,000$ evaluations of the periodogram. On the other hand, the LINEAR data shown in Figure 1 do not have any notable aliasing structure in their window function. In this case the maximum detectable frequency is the Nyquist limit defined by its temporal resolution, which is six digits beyond the decimal point in days. From Equation (30), we can write $f_{\text{Ny}} = 500,000 \text{ day}^{-1}$, and given the observing window of $T = 1962$

days, we find that five evaluations per peak across the entire detectable frequency range would require $N_{\text{eval}} \approx 4.9 \times 10^9$ evaluations of the periodogram! Computing this large a periodogram in most cases is computationally intractable (see Section 7.6), and so in practice one must choose a smaller limiting frequency based on prior information about what kind of signals are expected in the data: for example, in Figure 2, we chose a limiting frequency of $f_{\text{max}} = (1 \text{ hr})^{-1}$ based on typical oscillation periods expected for SX Phe-type stars. This leads to a much more manageable $N_{\text{eval}} \approx 240,000$ periodogram evaluations.

By comparison, data from the LSST survey (Ivezić et al. 2008) will fall somewhere in between: full frequency coverage of the 10 yr data up to a limiting frequency defined by the 30 s integration time for each visit would require roughly 25 million periodogram evaluations per object, which for fast implementations (see the next section) could be accomplished in several seconds on a modern CPU.

One final note: although it can be more easily interpretable to visualize periodograms as a function of period rather than a function of frequency, the peak widths are not constant in period. Regular grids in period rather than frequency tend to oversample at large periods and undersample at small periods; for this reason it is preferable to use a regular grid in frequency.

7.2. Failure Modes

When using the Lomb–Scargle periodogram in an observational pipeline, it is vital to keep in mind the failure modes of the periodogram approach, which are related to the aliasing and pseudo-aliasing effects rooted in the structure of the window function (recall Section 3). Owing to the interaction of the signal, the convolution due to the survey window, and noise in the data, it is quite common for the largest peak in the periodogram to correspond not to the true frequency but to some alias of that frequency.

Figure 24 demonstrates this for some simulated data. The data consist of 60 noisy observations each of 1000 simulated light curves within a span of 180 days. The window is typical of ground-based data, with each observation recorded within a few hours before or after midnight each night. The left panel shows the results: the periodogram peak coincides with the true period in under 50% of cases, and the modes of failure lead to noticeable patterns in the resulting plot. Though these are

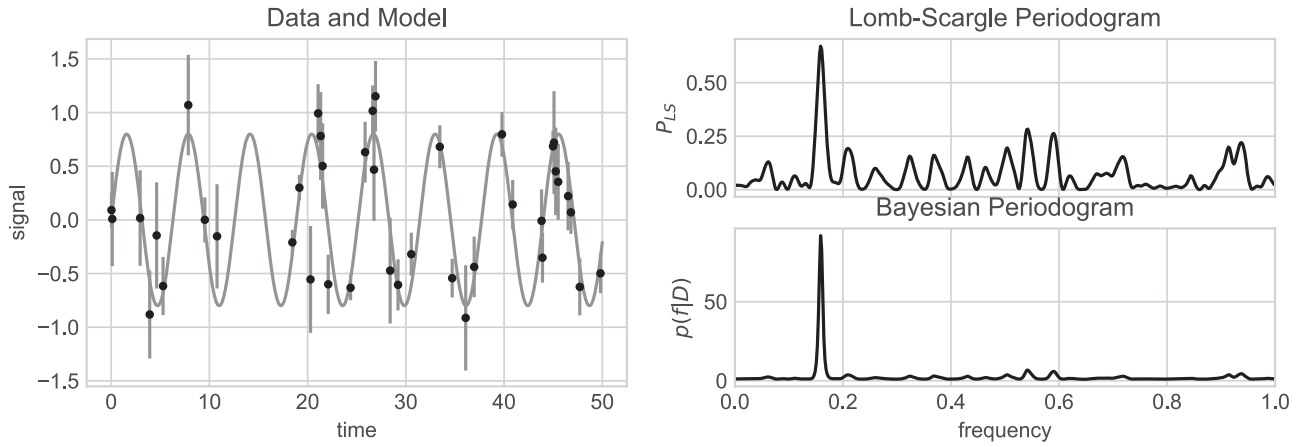


Figure 22. Comparison of the Lomb–Scargle periodogram (top right) and the Bayesian posterior periodogram (bottom right) for simulated data (left) drawn from a simple sinusoid. The Bayesian posterior is equal to the exponentiation of the unnormalized Lomb–Scargle periodogram and thus tends to suppress all but the largest peak. The Bayesian approach can be useful but is problematic if not used carefully (see text for discussion).

simulated observations, the pattern seen here is typical of real observations: see, for example, VanderPlas & Ivezić (2015) and Long et al. (2016), which show similar plots for RR Lyrae candidates from the Sloan Digital Sky Survey.

These patterns can be understood in terms of the interaction between the window function and the underlying spectral power. As we discussed in Section 3, a nightly observation pattern—typical of ground-based surveys—leads to a window function with a strong diurnal component that causes each frequency signature f_0 to be partially aliased at $f_0 + n\delta f$, for integers n and $\delta f = 1 \text{ cycle day}^{-1}$. (Recall Figures 13–14). This is the source of the failure modes depicted in panel (a) of Figure 24; in terms of periods the above expression becomes

$$P_{\text{obs}} = \left(\frac{1}{P_{\text{true}}} + n \delta f \right)^{-1}, \quad (45)$$

where, to be clear, n is a positive or negative integer and δf is the frequency of a strong feature in the window (here, 1 cycle day^{-1}). For the simulated data, close to 36% of the objects are mischaracterized along these failure modes.

Another mode of failure is when the periodogram for an object with frequency f_0 isolates a higher harmonic of that fundamental frequency, such as $2f_0$. This may occur for periodic signals that are not strictly sinusoidal and so have power at a higher harmonics. This higher harmonic peak is also subject to the same aliasing effects as in Equation (45). Thus, we can extend Equation (45) and describe the failure modes depicted in panel (b) of Figure 24 with the following equation, for positive integers $m > 0$:

$$P_{\text{obs}} = \left(\frac{m}{P_{\text{true}}} + n \delta f \right)^{-1}. \quad (46)$$

Panel (b) of Figure 24 illustrates this for $m = 2$; this harmonic and its aliases account for roughly 15% of the results.

Panel (c) of Figure 24 shows the final pattern of failure in Lomb–Scargle results, which has an opposite trend with true period. This is closely related to the effects shown in panels (a) and (b) but comes from the even symmetry of the Lomb–Scargle periodogram. Every periodogram peak at frequency f_0 has a corresponding peak at $-f_0$, and this is true of aliases as well. If a peak’s aliases cross into the negative frequency regime, they are effectively reflected into the positive-

frequency range. This reflection can be accounted for by a further modification of Equation (46)—taking its absolute value:

$$P_{\text{obs}} = \left| \frac{m}{P_{\text{true}}} + n \delta f \right|^{-1}. \quad (47)$$

Panel (c) shows the 5% of objects that fall along this reflected failure mode for $m = 1$, $n = -2$. After accounting for all these known sources of periodogram failure, only roughly 1% of points are misclassified in an “unexplained” way, seen in panel (d).

When applying a periodogram in practice, it is vital to take such effects into account, rather than blindly relying on the single periodogram peak as your best estimate of the period. Applying understanding of windowing and aliasing effects can help in detecting failures of the periodogram, but it is no silver bullet. For an observed peak at f_{peak} from a survey whose window has strong power at δf , something like the following should probably be employed:

1. Check for a peak at f_{peak}/m for at least $m \in \{2, 3\}$. If a significant peak is found, then f_{peak} is probably an order- m harmonic of the true frequency.⁶
2. Check for peaks at $|f_{\text{peak}} \pm n \delta f|$ for at least $n \in \{1, 2\}$ (where δf is determined from plotting the survey window power and is generally $(1 \text{ day})^{-1}$ for ground-based surveys). If these aliases exist, then it is possible that you have found a peak on the sequence of expected aliases—though keep in mind that there is no way to know from the periodogram alone whether or not this is the “true” peak!
3. For each of the top few of these aliases, fit a more complicated model (such as a multiterm Fourier series, template-based fit, etc.) to select among them.

For noisy observations, this procedure cannot generally guarantee that you have found the correct peak, but it is far preferable to the simplistic approach of blindly trusting the highest peak in the periodogram! We will come back to the question of uncertainty in the periodogram result in Section 7.4.

⁶ Notice that this step can detect aliases like those encountered in Figure 19, without resorting to a problematic multiterm model.




Figure 23. Example of a poorly chosen frequency grid for the data in Figure 1. The dark curve shows a periodogram evaluated on a grid of 10,000 equally spaced periods; the light curve shows the true periodogram (evaluated at $\sim 200,000$ equally spaced frequencies). This demonstrates that using too coarse a grid can lead to a periodogram that entirely misses relevant peaks.

7.3. Window Functions and Deconvolution

Our discussion of window functions here and in Section 3 has highlighted the impact of structure in the survey window on the resulting observed power spectrum and the importance of examining that window power to understand features of the resulting periodogram. Here we will consider the computation of the window function, as well as the possibility of recovering the true periodogram by deconvolving the window.

7.3.1. Computing the Window Function

The window power spectrum can be computed directly from the delta-function representation of the window function; From Equations (1), (9), and (27), we can write

$$\mathcal{P}_W(f; \{t_n\}) = \left| \sum_{n=1}^N e^{-2\pi i f t_n} \right|^2. \quad (48)$$

Comparing this to Equation (26), we see that this is essentially the classical periodogram for data $g_n = 1$ at all times t_n . With this fact in mind, one convenient way to estimate the form of the window power is to compute a standard Lomb–Scargle periodogram on a series of unit measurements, making sure *not* to pre-center the data or to use a floating-mean model. As Scargle (1982) notes, this computation does *not* give the true window—it differs from the true window just as the Lomb–Scargle periodogram differs from the classical periodogram—but in practice it is accurate enough to give useful insight into the window function’s important features. This method is how window power spectra have been computed throughout this paper.

7.3.2. Deconvolution and CLEANing

With this ability to compute the window function, we might hope to be able to use it *quantitatively* to remove spurious peaks from the observed power spectrum. Recall from Equation (28) that the observed data are a pointwise product of the underlying function and the survey window,

$$g_{\text{obs}}(t) = g(t)W_{\{t_n\}}(t), \quad (49)$$

and the convolution theorem tells us that the observed Fourier transform is a convolution of the true transform and the

window transform,

$$\mathcal{F}\{g_{\text{obs}}\} = \mathcal{F}\{g\} * \mathcal{F}\{W_{\{t_n\}}\}. \quad (50)$$

Given this relationship, we might hope to be able to invert this convolution to recover $\mathcal{F}\{g\}$ directly. For example, we could write

$$\begin{aligned} g(t) &= g_{\text{obs}}(t)/W_{\{t_n\}}(t) \\ \mathcal{F}\{g\} &= \mathcal{F}\{g_{\text{obs}}/W_{\{t_n\}}\} \\ &= \mathcal{F}\{g_{\text{obs}}\} * \mathcal{F}\{1/W_{\{t_n\}}\}. \end{aligned} \quad (51)$$

Because of the localization of observations, $W_{\{t_n\}}(t)$ is zero for most values of t , and so $1/W_{\{t_n\}}(t)$ and its Fourier transform are not well defined. Because of this, direct deconvolution is not an option in most cases of interest—in other words, the deconvolution problem is ill posed and does not have a unique solution.

There have been a few attempts in the literature to use procedural algorithms to solve this underconstrained deconvolution problem, perhaps most notably by adapting the iterative CLEAN algorithm developed for deconvolution in the context of radio astronomy (Roberts et al. 1987). For cleaning of Lomb–Scargle periodograms, the CLEAN approach is hindered by three main deficiencies: First and most importantly, the CLEAN algorithm at each iteration assumes that the highest peak is the location of the primary signal; this is not always borne out for realistic observations of faint objects where the cleaning is most necessary (recall the discussion in Section 7.2). Second, the convolution takes place at the level of the Fourier transform rather than the PSD; trying to clean a PSD directly ignores important phase information. Third, the CLEAN algorithm assumes a classical fast Fourier transform analysis: while the Lomb–Scargle periodogram is equivalent to a classical periodogram in the limit of equally spaced observations, it is not equivalent in the relevant case of unequal observations (see Section 5), and so any attempt to apply CLEAN to Lomb–Scargle analysis directly would be fundamentally flawed even if it were not ill posed to begin with.

The latter two issues could be remedied by focusing on the nonuniform Fourier transform and the classical periodogram rather than the Lomb–Scargle modification, but doing so would jettison the benefits of Lomb–Scargle—i.e., its useful statistical properties and extensibility via least squares—and this would still not address the far more problematic first issue. My feeling

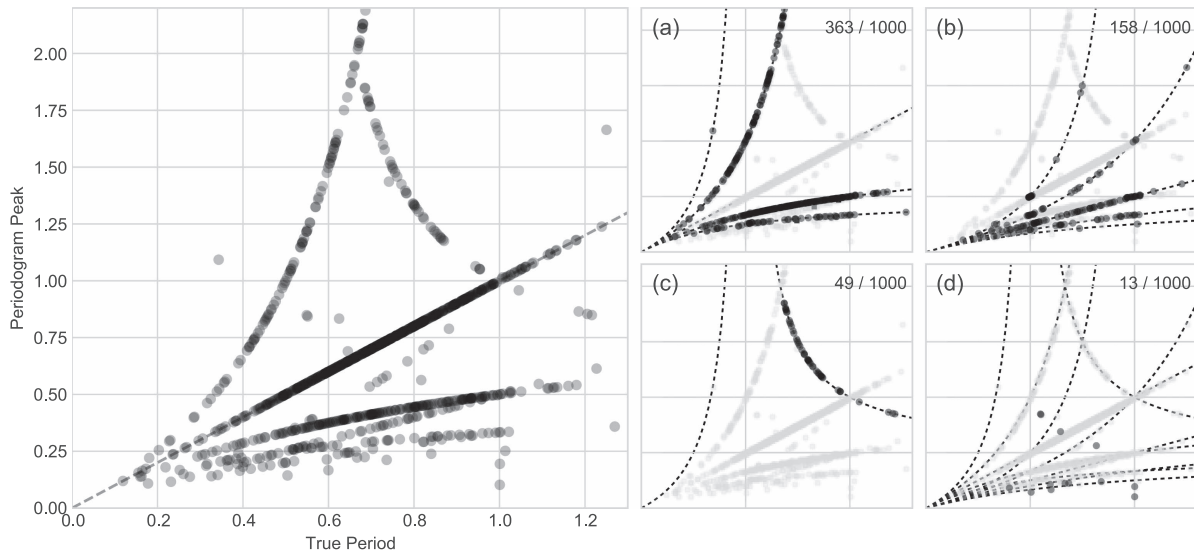


Figure 24. Comparison of the true period and peak Lomb–Scargle period for 1000 simulated periodic light curves. Each has 60 irregular observations over 180 nights, with a cadence typical of ground-based surveys (i.e., showing a strong diurnal window pattern similar to Figure 13). The Lomb–Scargle peak does not always coincide with the true period, and there is noticeable structure among these failures. Panels (a)–(d) isolate some of the specific modes of failure that should be expected for this kind of window function; see the text for more discussion.

is that there is room for a more principled approach to the unconstrained deconvolution problem for periodograms derived from nonuniform fast Fourier transforms (perhaps through some sort of L1/lasso regularization that imposes assumptions of sparsity on the true periodogram), but to date this does not appear to have been explored anywhere in the literature.

7.4. Uncertainties in Periodogram Results

An important aspect of reporting results from the Lomb–Scargle periodogram is the uncertainty of the estimated period. In many areas of science we are used to being able to report uncertainties in terms of error bars, e.g., “the period is 16.3 ± 0.6 hr.” For periods derived from the Lomb–Scargle periodogram, however, uncertainties generally cannot be meaningfully expressed in this way: as we saw in the discussion of failure modes in Section 7.2, the concern for periodograms is more often a disjointed inaccuracy associated with false peaks and aliases, rather than a more smooth imprecision in location of a particular peak.

7.4.1. Peak Width and Frequency Precision

A periodic signal will be reflected in the periodogram by the presence of a peak of a certain width and height. In the Fourier view, the precision with which a peak’s frequency can be identified is directly related to the width of this peak; often the half-width at half-maximum $f_{1/2} \approx 1/T$ is used. This can be formalized more precisely in the least-squares interpretation of the periodogram, in which the inverse of the curvature of the peak is identified with the uncertainty (Ivezić et al. 2014)—which in the Bayesian view amounts to fitting a Gaussian curve to the (exponentiated) peak (Jaynes 1987; Bretthorst 1988). This introduces first-order dependence on the number of samples N and their average signal-to-noise ratio Σ ; the scaling

is approximately (see, e.g., Gregory 2001)

$$\sigma_f \approx f_{1/2} \sqrt{\frac{2}{N\Sigma^2}}. \quad (52)$$

This dependence comes from the fact that the Bayesian uncertainty is related to the width of the *exponentiated* periodogram, which depends on P_{\max} , the height of the peak.⁷

We have motivated from the Fourier window discussions that peak width $f_{1/2}$ is the inverse of the observational baseline; the surprising result is that to first order the peak width in the periodogram *does not depend* on either the number of observations or their signal-to-noise ratio! This can be seen visually in Figure 25, which shows periodograms for simulated data with a fixed four-period baseline, with varying sample sizes and signal-to-noise ratios. In all cases, the widths of the primary peak are essentially *identical*, regardless of the quality or quantity of data! Data quality and quantity are reflected in the height of the peak in relation to the “background noise,” which speaks to the peak *significance* rather than the precision of frequency detection.

For this reason, if you insist on reporting frequency error bars derived from peak width, the results can be pretty silly for observations with long baselines. For example, going by the peak width in Figure 2, the periodogram reveals a period of 2.58014 ± 0.00004 hr, a relative precision of about 1/1000 percent. While this is an accurate characterization of the period precision *assuming that peak is the correct one*, it does not capture the much more relevant uncertainties demonstrated in Section 7.2, in which we might find ourselves on the wrong peak entirely. This is why peak widths and Gaussian error bars

⁷ To see why, consider a periodogram with maximum value $P_{\max} = P(f_{\max})$, so that $P(f_{\max} \pm f_{1/2}) = P_{\max}/2$. The Bayesian uncertainty comes from approximating the exponentiated peak as a Gaussian, i.e., $\exp[P(f_{\max} \pm \delta f)] \propto \exp[-\delta f^2/(2\sigma_f^2)]$. From this we can write $P_{\max}/2 \approx P_{\max} - f_{1/2}^2/(2\sigma_f^2)$ or $\sigma_f \approx f_{1/2}/\sqrt{P_{\max}}$. In terms of signal-to-noise ratio $\Sigma = \text{rms}[(y_n - \mu)/\sigma_n]$, a well-fit model gives $P_{\max} \approx \chi_0^2/2 \approx \Sigma^2 N/2$, which leads to the expression in Equation (52).




Figure 25. Effect of the number of points N and the signal-to-noise ratio (S/N) on the expected width and height of the periodogram peak. Top panels show the normalized periodogram (Equation (59)), while bottom panels show the PSD-normalized periodogram (Equation (33)) scaled by noise variance. Perhaps surprisingly, neither the number of points nor the signal-to-noise ratio affects the peak width.

should generally be avoided when reporting uncertainties in the context of a periodogram analysis.

7.4.2. False-alarm Probability

A much more relevant quantity for expressing uncertainty of periodogram results is the height of the peak, and in particular the height compared to the spurious background peaks that arise in the periodogram. Figure 25 indicates that this property *does* depend on both the number of observations and their signal-to-noise ratio: for the small- N /low signal-to-noise ratio cases, the spurious peaks in the background become comparable to the size of the true peak. In fact, as we saw in Section 5, the ability to analytically define and quantify the relationship between peak height and significance is one of the primary considerations that led to the standard form of the Lomb-Scargle periodogram.

The typical approach to quantifying the significance of a peak is the FAP, which measures the probability that a data set with no signal would—owing to coincidental alignment among the random errors—lead to a peak of a similar magnitude. Scargle (1982) showed that for data consisting of pure Gaussian noise, the values of the unnormalized periodogram in Equation (33) follow a χ^2 distribution with two degrees of freedom, that is, at any given frequency f_0 , if $Z = P(f_0)$ is the periodogram value from Equation (33), then

$$P_{\text{single}}(Z) = 1 - \exp(-Z) \quad (53)$$

is the cumulative probability of observing a periodogram value less than Z , in data consisting only of Gaussian noise.⁸

7.4.2.1. Independent Frequency Method

We are generally not interested in the distribution of one particular randomly chosen frequency, but rather the distribution of *the highest peak* of the periodogram, which is a quite different situation.

By analogy, consider rolling a standard six-sided die. The probability, in a single roll, of rolling a number, say, $r > 4$ is easy to compute: it is two sides out of six, or $p(r > 4) = 1/3$. If, on the

other hand, you roll 10 dice and choose the largest number among them, the probability that it will be greater than 4 is *far* larger than $1/3$; it is $p(\max_{10}(r) > 4) = 1 - (1 - 1/3)^{10} \approx 0.98$. The case for the periodogram is analogous: the probability of seeing a spurious peak at any single location (Equation (53)) is relatively small, but the probability of seeing a single spurious peak among a large number of frequencies is much higher.

With the dice, this calculation is easy because the rolls are independent: the result of one roll does not affect the result of the next. With the periodogram, though, the value at one frequency is correlated with the value at other frequencies in a way that is quite difficult to analytically express—these correlations come from the convolution with the survey window. Nonetheless, one common approach to estimating the distribution has been to assume that it can be modeled on some “effective number” of independent frequencies N_{eff} , so that the FAP can be estimated as

$$FAP(z) \approx 1 - [P_{\text{single}}(z)]^{N_{\text{eff}}}. \quad (54)$$

A very simple estimate for N_{eff} is based on our arguments of the expected peak width, $\delta f = 1/T$. In this approximation, the number of independent peaks in a range $0 \leq f \leq f_{\text{max}}$ is assumed to be $N_{\text{eff}} = f_{\text{max}}T$. There have been various attempts to estimate this value more carefully, both analytically and via simulations (see, e.g., Horne & Baliunas 1986; Schwarzenberg-Czerny 1998; Cumming 2004; Frescura et al. 2008), but all such approaches are necessarily only approximations.

7.4.2.2. Baluev (2008) Method

A more sophisticated treatment of the problem is that of Baluev (2008), who derived an analytic result based on the theory of extreme values for stochastic processes. For the standard periodogram in Equation (33), Baluev (2008) showed that the following formula for the FAP provides a close upper limit even in the case of highly structured window functions:

$$FAP(z) \approx 1 - P_{\text{single}}(z)e^{-\tau(z)}, \quad (55)$$

where, for the normalized periodogram (Equation (59)),

$$\tau(z) \approx W(1 - z)^{(N-4)/2} \sqrt{z} \quad (56)$$

and $W = f_{\text{max}} \sqrt{4\pi \text{var}(t)}$ is an effective width of the observing window in units of the maximum frequency chosen for the

⁸ Be aware that for different periodogram normalizations (see Section 7.5) the form of this distribution changes; see Cumming et al. (1999) or Baluev (2008) for a good summary of the statistical properties of various periodogram normalizations.




Figure 26. Comparison of various approaches to computing the FAP for simulated observations with both structured and unstructured survey windows. Note that the approximate N_{eff} method in its most naive form tends to underestimate the bootstrapped FAP, while the Baluev (2008) method tends to overestimate the bootstrapped FAP, particularly for data with a highly structured survey window.

analysis. This result is not an exact measure of the FAP, but rather an upper limit that is valid for alias-free periodograms, i.e., cases where the window function has very little structure, but which holds quite well even in the case of more realistic survey windows.

7.4.2.3. Bootstrap Method

In the absence of a true analytic solution to the FAP, we can turn to computational methods such as the bootstrap. The bootstrap method is a technique in which the statistic in question is computed repeatedly on many random resamplings of the data in order to approximate the distribution of that statistic (see Ivezić et al. 2014, for a useful general discussion of this technique). For the periodogram, in each resampling we keep the temporal coordinates the same, draw observations randomly with replacement from the observed values, and then compute the maximum of the resulting periodogram. For enough resamplings, the distribution of these maxima will approximate the true distribution for the case with no periodic signal present. The bootstrap produces the most robust estimate of the FAP because it makes few assumptions about the form of the periodogram distribution and fully accounts for survey window effects.

Unfortunately, the computational costs of the bootstrap can be quite prohibitive: to accurately measure small levels of FAP requires constructing a large number of full periodograms. If you are probing a false-positive rate of r among N bootstrap resamples, you would roughly expect to find $rN \pm \sqrt{rN}$ relevant peaks in your bootstrap sample. More concretely, for 1000 bootstrap samples, you will find only $\sim 10 \pm 3$ peaks reflecting a 1% false-positive rate. The random fluctuations in that count translate directly to an inability to accurately estimate the false-positive rate at that level. A good rule of thumb is that to accurately characterize a false-positive rate r requires something like $\sim 10/r$ individual periodograms to be computed—this can

grow computationally intractable quite quickly. The bootstrap is also not universally applicable: for example, it does not correctly account for cases where noise in observations is correlated; for more discussion of the bootstrap approach and its caveats, see Ivezić et al. (2014) and references therein.

Figure 26 compares these methods of estimating the FAP, for both an unstructured survey window and a structured window that produces the kinds of aliasing we have discussed above. The naive approximation in this case underestimates the FAP at nearly all levels—so, for example, it might lead you to think that a peak has an FAP of 10% when in fact it is closer to 30%. The Baluev method, by design, overestimates the FAP and is quite close to the bootstrapped distribution in the case of an unstructured window. For a highly structured window, the Baluev method does not perform as well, but it still tends to overestimate the FAP—so, for example, it might lead you to think of a peak as an FAP of 10% when in fact it is closer to 5%.

Finally, I will note that Suveges (2012) has shown the promise of a hybrid approach of the bootstrap method and the extreme value statistics of Baluev (2008); this has the ability to compute accurate FAPs without the need to compute thousands of full bootstrapped periodograms.

7.4.2.4. Detection Significance with Red Noise

It is important to recognize that the above methods of measuring peak significance or FAP generally depend on the assumption of white noise: that is, that the height of spurious “background” peaks in the periodogram is independent of frequency. In many situations of interest, however, this assumption does not hold: for example, compact accreting objects often display “red noise”—that is, noise that increases in power toward lower frequencies (Vaughan et al. 2003). The presence of red noise complicates both the detection of peaks and the estimation of their significance or FAP; for detailed

discussion of various approaches to periodic analysis in this context, see Vaughan (2005, 2010).

7.4.2.5. What Does the False-alarm Probability Measure?

While the FAP can be a useful concept, you should be careful to keep in mind that it answers a very specific question: “What is the probability that a signal with no periodic component would lead to a peak of this magnitude?” In particular, it emphatically does *not* answer the much more relevant question, “What is the probability that this data set is periodic given these observations?” This boils down to a case of $P(A|B) \neq P(B|A)$, but still it is common to see the FAP treated as if it speaks to the second rather than the first question.

Similarly, the FAP tells us nothing about the false-negative rate, i.e., the rate at which we would expect to miss a periodic signal that is, in fact, present. It also tells us nothing about the error rate, i.e., the rate at which we would expect to identify an incorrect alias of a signal present in our data (as we saw in Section 7.2). Too often users are tempted to interpret the FAP too broadly, with the hope that it would speak to questions that are beyond its reach.

Unfortunately, there is no single, universally applicable answer to these broader, more relevant questions of uncertainty of Lomb–Scargle results. Perhaps the most fruitful path toward understanding of such effects for a particular set of observations—with particular noise characteristics and a particular observing window—is via simulated data injected into the detection pipeline. In fact, this type of simulation is a vital component of most (if not all) current and future time-domain surveys that seek to detect, characterize, and catalog periodic objects, regardless of the particular approach used to identify periodicity (see, e.g., Delgado et al. 2006; Ivezić et al. 2008; Sesar et al. 2010; McQuillan et al. 2012; Oluseyi et al. 2012; Ridgway et al. 2012; Drake et al. 2014, etc.)

7.5. Periodogram Normalization and Interpretation

Often it is useful to be able to interpret the periodogram values themselves. Recall that there are two equally valid perspectives of what the periodogram is measuring—the Fourier view and the least-squares view—and each of these lends itself to a different approach to normalization and interpretation of the periodogram.

7.5.1. PSD Normalization

When considering the periodogram from the Fourier perspective, it is useful to normalize the periodogram such that in the special case of equally spaced data it recovers the standard Fourier power spectrum. This is the normalization used in Equation (33) and the equivalent least-squares expression seen in Equation (37):

$$P(f) = \frac{1}{2} [\hat{\chi}_0^2 - \hat{\chi}^2(f)]. \quad (57)$$

For equally spaced data, this periodogram is identical to the standard definition based on the fast Fourier transform:

$$P(f) = \frac{1}{N} |\text{FFT}(y_n)|^2; \quad (58)$$

in particular, this means that the units of the periodogram are $\text{unit}(y)^2$ and can be interpreted as squared amplitudes of the Fourier component at each frequency. Note, however, that the

units change if data uncertainties are incorporated into the periodogram as in Section 6.1; the periodogram in this normalization becomes unitless: it is essentially a measure of periodic content in signal-to-noise ratio rather than in signal itself.

7.5.2. Least-squares Normalization

In the least-squares view of the periodogram, the periodogram is interpreted as an inverse measure of the goodness of fit for a model. When we express the periodogram as a function of χ^2 model fits as in Equation (57), it becomes clear that the periodogram has a mathematical maximum value: if the sinusoidal model *perfectly* fits the data at some frequency f_0 , then $\hat{\chi}^2(f_0) = 0$ and the periodogram is maximized at a value of $\hat{\chi}_0^2/2$. On the other end, it is mathematically impossible for a best-fit sinusoidal model at any frequency to do worse than the simple constant, nonvarying fit reflected in $\hat{\chi}_0^2$, and so the minimum value of Equation (57) is exactly zero.

This suggests a different normalization of the periodogram, where the values are dimensionless and lie between 0 and 1:

$$P_{\text{norm}}(f) = 1 - \frac{\hat{\chi}^2(f)}{\hat{\chi}_0^2}. \quad (59)$$

This “normalized periodogram” is unitless and is directly proportional to the unnormalized (or PSD-normalized) periodogram in Equation (57).

While the normalization does not affect the shape of the periodogram, its main practical consequence is that the statistics of the resulting periodogram differ slightly, and this needs to be taken into account when computing analytic estimates of uncertainties and FAPs explored in Section 7.4.2. Other normalizations exist as well, but they seem to be rarely used in practice. For a concise summary of several periodogram normalizations and their statistical properties, refer to the introduction of Baluev (2008).

7.6. Algorithmic Considerations

Given the size of the frequency grid required to fully characterize the periods from a given data set, it is vital to have an efficient algorithm for evaluating the periodogram. We will see that the naive implementation of the standard Lomb–Scargle formula (Equation (33)) scales poorly with the size of the data, but that fast alternatives are available.

Suppose you have a set of N observations over a time span T for an average cadence of $\bar{\delta t} = N/T$. From Equation (44) we see that the number of frequencies we need to evaluate is $N_f \propto T f_{\text{max}} = N f_{\text{max}} / \bar{\delta t}$. Holding constant the average survey cadence $\bar{\delta t}$ and f_{max} , we find that the number of frequencies required is directly proportional to the number of data points. The computation of the Lomb–Scargle periodogram in Equation (33) requires sums over N sinusoids at each of the N_f frequencies, and thus we see that the naive scaling of the algorithm with the size of the data set is $\mathcal{O}(N^2)$, when survey properties are held constant. Due to the trigonometric functions involved, this turns out to be a rather “expensive” $\mathcal{O}(N^2)$, which makes the direct implementation impractical for even modestly sized data sets.

Fortunately, several faster implementations have been proposed to compute the periodogram to arbitrary precision in $\mathcal{O}(N \log N)$ time. The first of these is discussed in Press &

Rybicki (1989), which uses an inverse interpolation operation along with a fast Fourier transform to compute the trigonometric components of Equation (33) very efficiently over a large number of frequencies. Zechmeister & Kürster (2009) showed how this approach can be straightforwardly extended to the floating-mean periodogram discussed in Section 6.2.

A qualitatively similar approach to the Press & Rybicki (1989) algorithm is presented by Leroy (2012): it makes use of the non-equispaced fast Fourier transform (NFFT; see Keiner et al. 2009) to compute the Lomb–Scargle periodogram about a factor of 10 faster than the Press & Rybicki (1989) approach. For the multiterm models discussed in Section 6.3, Palmer (2009) presents an adaptation of the Press & Rybicki (1989) method that can compute the multiterm result in $\mathcal{O}(NK \log N)$, for N data points and K Fourier components.

8. Conclusion and Summary

This paper has been a conceptual tour of the Lomb–Scargle periodogram, from its roots in Fourier analysis to its equivalence with special cases of periodic analysis based on least-squares model fitting and Bayesian analysis. From this conceptual understanding, we considered a list of challenges and issues to be considered when applying the method in practice. We will finish here with a brief summary of these practical recommendations, along with a somewhat opinionated postscript for further thought.

8.1. Summary of Recommendations

The previous pages contain a large amount of background and advice for working with the Lomb–Scargle periodogram. Following is a brief summary of the considerations to keep in mind when you apply this algorithm to a data set:

1. Choose an appropriate frequency grid: the minimum can be set to zero, the maximum set based on the precision of the time measurements (Section 4.1.2), and the grid spacing set based on the temporal baseline of the data (Section 7.1) so as not to sample too coarsely around peaks. If this grid size is computationally intractable, reduce the maximum frequency based on what kinds of signals you are looking for.
2. Compute the window transform using the Lomb–Scargle periodogram, by substituting $g_n = 1$ for each t_n and making sure not to pre-center the data or use a floating-mean model (Section 7.3.1). Examine this window function for dominant features, such as daily or annual aliases (see Figure 13) or Nyquist-like limits (see Figure 15).
3. Compute the periodogram for your data. You should always use the floating-mean model (Section 6.2), as it produces more robust periodograms and has few if any disadvantages. Avoid multiterm Fourier extensions (Section 6.3) when the signal has an unknown form, because its main effect is to increase periodogram noise (see Figures 20–21).
4. Plot the periodogram and identify any patterns that may be caused by features you observed in the window function power. Plot reference lines showing several FAP levels to understand whether your periodogram peaks are significant enough to be labeled detections: use the Baluev method or the bootstrap method if it is computationally feasible (Section 7.4.2). Keep in mind

exactly what the FAP measures and avoid the temptation to misinterpret it (Section 7.4.2.5).

5. If the window function shows strong aliasing, locate the expected multiple maxima and plot the phased light curve at each. If there is indication that the sinusoidal model underfits the data (see Figure 19), then consider refitting with a multiterm Fourier model (Section 6.3).
6. If you have prior knowledge of the shape of light curves you are trying to detect, consider using more complex models, such as multiterm models or physically derived templates, to choose between multiple peaks in the periodogram (Section 6.5). This type of refinement can be quite useful in building automated pipelines for period fitting, especially in cases where the window aliasing is strong.
7. If you are building an automated pipeline based on Lomb–Scargle for use in a survey, consider injecting known signals into the pipeline to measure your detection efficiency as a function of object type, brightness, and other relevant characteristics (Section 7.2)

This list is certainly not comprehensive for all uses of the periodogram, but it should serve as a brief reminder of the kinds of issues you should keep in mind when using the method to detect periodic signals.

8.2. Postscript: Why Lomb–Scargle?

After considering all of these practical aspects of the periodogram, I think it is worth stepping back to revisit the question of *why* astronomers tend to gravitate toward the Lomb–Scargle approach rather than the (in many ways simpler) classical periodogram.

As discussed in Section 5, the Lomb–Scargle approach has two distinct benefits over the classical periodogram: the noise distribution at each individual frequency is chi-square distributed under the null hypothesis, and the result is equivalent to a periodogram derived from a least-squares analysis. But somehow along the way a mythology seems to have developed surrounding the efficiency and efficacy of the Lomb–Scargle approach. In particular, it is common to see statements or implications along the lines of the following:

1. *Myth: the Lomb–Scargle periodogram can be computed more efficiently than the classical periodogram.* Reality: computationally, the two are quite similar, and in fact the fastest Lomb–Scargle algorithm currently available is based on the classical periodogram computed via the NFFT algorithm (see Section 7.6).
2. *Myth: the Lomb–Scargle periodogram is faster than a direct least-squares periodogram because it avoids explicitly solving for model coefficients.* Reality: model coefficients can be determined with little extra computation (see the discussion in Ivezić et al. 2014).
3. *Myth: the Lomb–Scargle periodogram allows analytic computation of statistics for periodogram peaks.* Reality: while this is true at individual frequencies, it is not true for the more relevant question of maximum peak heights across multiple frequencies, which must be either approximated or computed by bootstrap resampling (see Section 7.4)
4. *Myth: the Lomb–Scargle periodogram corrects for aliasing due to sampling and leads to independent noise at each frequency.* Reality: for structured window

functions common to most astronomical applications, the Lomb–Scargle periodogram has the same window-driven issues as the classical periodogram, including spurious peaks due to partial aliasing and highly correlated periodogram errors (see Section 7.2).

5. *Myth: Bayesian analysis shows that Lomb–Scargle is the optimal statistic for detecting periodic signals in data.* Reality: Bayesian analysis shows that Lomb–Scargle is the optimal statistic for *fitting a sinusoid to data*, which is not the same as saying that it is optimal for finding the frequency of a generic, potentially nonsinusoidal signal (see Section 6.5).

With these misconceptions corrected, what is the practical advantage of Lomb–Scargle over a classical periodogram? What would we lose if we instead used the simple classical Fourier periodogram, estimating uncertainty, significance, and FAP by resampling and simulation, as we must for Lomb–Scargle itself?

The advantage of analytic statistics for Lomb–Scargle evaporates in light of the need to account for multiple frequencies, so the only advantage left is the correspondence to least-squares and Bayesian models, and in particular the ability to generalize to more complicated models where appropriate—but in this case you are not really using Lomb–Scargle at all, but rather a generative Bayesian model for your data based on some strong prior information about the form of your signal. The equivalence of Lomb–Scargle to a Bayesian sinusoidal model is perhaps an interesting bit of trivia, but not itself a reason to use that model if your data are not known a priori to be sinusoidal—it could even be construed as an argument *against* Lomb–Scargle in the general case where the assumption of a sinusoid is not well-founded.

Conversely, if you replace your Lomb–Scargle approach with a classical periodogram, what you gain is the ability to reason quantitatively about the effects of the survey window function on the resulting periodogram (see Section 7.3.2). While the deconvolution problem is ill-posed, there is no reason to assume that this is a fatal defect: ill-posed linear models are solved routinely in other areas of computational research, particularly by using sparsity priors or various forms of regularization. In any case, I would contend that there is ample room for practitioners to question the prevailing folk wisdom of the advantage of Lomb–Scargle over approaches based directly on the Fourier transform and classical periodogram.

8.3. Figures and Code

All computations and figures in this paper were produced using Python, and in particular used the Numpy (van der Walt et al. 2011; Oliphant 2015), Pandas (McKinney 2010), AstroPy (Astropy Collaboration et al. 2013), and Matplotlib (Hunter 2007) packages. Periodograms were computed using the AstroPy implementations⁹ of the Press & Rybicki (1989), Zechmeister & Kürster (2009), and Palmer (2009) algorithms, which were adapted from the Python code originally published by Ivezić et al. (2014) and VanderPlas & Ivezić (2015). All code behind this paper, including code to reproduce all figures, is available in the form of Jupyter notebooks in the paper’s GitHub repository.¹⁰

⁹ The AstroPy Lomb–Scargle documentation is at <http://docs.astropy.org/en/stable/stats/lombscargle.html>.

¹⁰ Source code and notebooks available at <http://github.com/jakevdp/PracticalLombScargle/>.

I am grateful to Jeff Scargle, Max Mähle, Dan Foreman-Mackey, Zeljko Ivezić, David Hogg, and the anonymous reviewer for helpful feedback on early drafts of this paper. This work was supported by the University of Washington eScience Institute, with funding from the Gordon and Betty Moore Foundation, the Alfred P. Sloan Foundation, and the Washington Research Foundation.

ORCID iDs

Jacob T. VanderPlas  <https://orcid.org/0000-0002-9623-3401>

References

- Anderson, T. 1971, *The Statistical Analysis of Time-Series*, Wiley Series in Probability and Mathematical Statistics: Probability and Mathematical Statistics (New York: Wiley)
- Angus, R., Foreman-Mackey, D., & Johnson, J. A. 2016, *ApJ*, **818**, 109
- Astropy Collaboration, Robitaille, T. P., Tollerud, E. J., et al. 2013, *A&A*, **558**, A33
- Baluev, R. V. 2008, *MNRAS*, **385**, 1279
- Brethorst, G. 1988, *Bayesian Spectrum Analysis and Parameter Estimation*, Lecture Notes in Statistics (Berlin: Springer)
- Brethorst, G. L. 2001, in AIP Conf. Ser. 568, *Bayesian Inference and Maximum Entropy Methods in Science and Engineering*, ed. A. Mohammad-Djafari (Melville, NY: AIP), 241
- Brethorst, G. L. 2003, in *Frequency Estimation and Generalized Lomb–Scargle Periodograms*, ed. E. D. Feigelson & G. J. Babu, 309–29
- Cumming, A. 2004, *MNRAS*, **354**, 1165
- Cumming, A., Marcy, G. W., & Butler, R. P. 1999, *ApJ*, **526**, 890
- Debosscher, J., Sarro, L. M., Aerts, C., et al. 2007, *A&A*, **475**, 1159
- Deeming, T. J. 1975, *Ap&SS*, **36**, 137
- Delgado, F., Cook, K., Miller, M., Allsman, R., & Pierfederici, F. 2006, *Proc. SPIE*, **6270**, 1
- Drake, A. J., Graham, M. J., Djorgovski, S. G., et al. 2014, *ApJS*, **213**, 9
- Dworetzky, M. M. 1983, *MNRAS*, **203**, 917
- Edelson, R. A., & Krolik, J. H. 1988, *ApJ*, **333**, 646
- Eyer, L., & Bartholdi, P. 1999, *A&AS*, **135**, 1
- Ferraz-Mello, S. 1981, *AJ*, **86**, 619
- Foster, G. 1996, *AJ*, **112**, 1709
- Frescura, F. A. M., Engelbrecht, C. A., & Frank, B. S. 2008, *MNRAS*, **388**, 1693
- Gilliland, R. L., & Baliunas, S. L. 1987, *ApJ*, **314**, 766
- Gottlieb, E. W., Wright, E. L., & Liller, W. 1975, *ApJL*, **195**, L33
- Graham, M. J., Drake, A. J., Djorgovski, S. G., et al. 2013a, *MNRAS*, **434**, 3423
- Graham, M. J., Drake, A. J., Djorgovski, S. G., Mahabal, A. A., & Donalek, C. 2013b, *MNRAS*, **434**, 2629
- Gregory, P. C. 2001, in AIP Conf. Ser. 568, *Bayesian Inference and Maximum Entropy Methods in Science and Engineering*, ed. A. Mohammad-Djafari (Melville, NY: AIP), 557
- Gregory, P. C., & Loredo, T. J. 1992, *ApJ*, **398**, 146
- Hilditch, R. W. 2001, in *An Introduction to Close Binary Stars* (Cambridge: Cambridge Univ. Press), 392
- Home, J. H., & Baliunas, S. L. 1986, *ApJ*, **302**, 757
- Huijse, P., Estevez, P. A., Protopapas, P., Zegers, P., & Principe, J. C. 2012, *ITSP*, **60**, 5135
- Huijse, P., Estevez, P. A., Zegers, P., Principe, J. C., & Protopapas, P. 2011, *ISPL*, **18**, 371
- Hunter, J. D. 2007, *CSE*, **9**, 90
- Irwin, A. W., Campbell, B., Morbey, C. L., Walker, G. A. H., & Yang, S. 1989, *PASP*, **101**, 147
- Ivezić, Ž., Connolly, A., VanderPlas, J., & Gray, A. 2014, *Statistics, Data Mining, and Machine Learning in Astronomy: A Practical Python Guide for the Analysis of Survey Data* (Princeton, NJ: Princeton Univ. Press)
- Ivezić, Ž., et al. 2008, arXiv:0805.2366
- Jaynes, E. T., & Brethorst, G. L. (ed.) 2003, *Probability Theory: The Logic of Science* (Cambridge: Cambridge Univ. Press)
- Jaynes, E. 1987, in *Maximum-Entropy and Bayesian Spectral Analysis and Estimation Problems*, ed. C. Smith & G. Erickson (The Netherlands: Springer), 1
- Keiner, J., Kunis, S., & Potts, D. 2009, *ACM Trans. Math. Softw.*, **36**, 19
- Kelly, B. C., Becker, A. C., Sobolewska, M., Siemiginowska, A., & Uttley, P. 2014, *ApJ*, **788**, 33
- Koen, C. 2006, *MNRAS*, **371**, 1390

- Kolenberg, K., Szabó, R., Kurtz, D. W., et al. 2010, [ApJL](#), **713**, L198
- Leroy, B. 2012, [A&A](#), **545**, A50
- Lomb, N. R. 1976, [Ap&SS](#), **39**, 447
- Long, J. P., Chi, E. C., & Baraniuk, R. G. 2016, [AnApS](#), **10**, 165
- McKinney, W. 2010, in Proc. 9th Python in Science Conf., Data Structures for Statistical Computing in Python, ed. S. van der Walt & J. Millman, 51
- McQuillan, A., Aigrain, S., & Roberts, S. 2012, [A&A](#), **539**, A137
- Mortier, A., Faria, J. P., Correia, C. M., Santerne, A., & Santos, N. C. 2015, [A&A](#), **573**, A101
- Oliphant, T. E. 2015, Guide to NumPy (2nd ed.; USA: CreateSpace Independent Publishing Platform)
- Oluseyi, H. M., Becker, A. C., Culliton, C., et al. 2012, [AJ](#), **144**, 9
- Palaversa, L., Ivezić, Ž., Eyer, L., et al. 2013, [AJ](#), **146**, 101
- Palmer, D. M. 2009, [ApJ](#), **695**, 496
- Percy, J. R. (ed.) 1986, The Study of Variable Stars Using Small Telescopes (Cambridge: Cambridge Univ. Press)
- Press, W. H., & Rybicki, G. B. 1989, [ApJ](#), **338**, 277
- Press, W. H., Teukolsky, S. A., Vetterling, W. T., & Flannery, B. P. 2007, Numerical Recipes 3rd Edition: The Art of Scientific Computing (3rd ed; New York: Cambridge Univ. Press)
- Reimann, J. D. 1994, PhD thesis, Univ. California, Berkeley
- Richards, J. W., Starr, D. L., Miller, A. A., et al. 2012, [ApJS](#), **203**, 32
- Ridgway, S. T., Chandrasekharan, S., Cook, K. H., et al. 2012, [Proc. SPIE](#), **8448**, 10
- Roberts, D. H., Lehar, J., & Dreher, J. W. 1987, [AJ](#), **93**, 968
- Scargle, J. D. 1982, [ApJ](#), **263**, 835
- Scargle, J. D. 1989, [ApJ](#), **343**, 874
- Scargle, J. D. 1998, [ApJ](#), **504**, 405
- Scargle, J. D. 2002, in AIP Conf. Ser. 617, Bayesian Inference and Maximum Entropy Methods in Science and Engineering, ed. R. L. Fry (Melville, NY: AIP), **23**
- Schuster, A. 1898, [TeMag](#), **3**, 13
- Schwarzenberg-Czerny, A. 1989, [MNRAS](#), **241**, 153
- Schwarzenberg-Czerny, A. 1996, [ApJL](#), **460**, L107
- Schwarzenberg-Czerny, A. 1998, [MNRAS](#), **301**, 831
- Schwarzenberg-Czerny, A. 1999, [ApJ](#), **516**, 315
- Sesar, B., Ivezić, Ž., Grammer, S. H., et al. 2010, [ApJ](#), **708**, 717
- Sesar, B., Grillmair, C. J., Cohen, J. G., et al. 2013, [ApJ](#), **776**, 26
- Sesar, B., Stuart, J. S., Ivezić, Ž., et al. 2011, [AJ](#), **142**, 190
- Stellingwerf, R. F. 1978, [ApJ](#), **224**, 953
- Suveges, M. 2012, in Seventh Conf. on Astronomical Data Analysis, ed. J.-L. Starck & C. Surace, **16**
- Swingler, D. N. 1989, [AJ](#), **97**, 280
- van der Walt, S., Colbert, S. C., & Varoquaux, G. 2011, [CSE](#), **13**, 22
- VanderPlas, J. T., & Ivezić, Z. 2015, arXiv:1502.01344
- Vaughan, S. 2005, [A&A](#), **431**, 391
- Vaughan, S. 2010, [MNRAS](#), **402**, 307
- Vaughan, S., Edelson, R., Warwick, R. S., & Uttley, P. 2003, [MNRAS](#), **345**, 1271
- Vetterli, M., Kovačević, J., & Goyal, V. K. 2014, Foundations of Signal Processing (Cambridge: Cambridge Univ. Press)
- Vio, R., Andreani, P., & Biggs, A. 2010, [A&A](#), **519**, A85
- Wang, Y., Khardon, R., & Protopapas, P. 2012, [ApJ](#), **756**, 67
- Zechmeister, M., & Kürster, M. 2009, [A&A](#), **496**, 577